

Week 4 - SQL: The Query Language

"Life is just a bowl of queries."

-Anon

Lecturer: Dr. Tony Stockman

Room CS405



Lecture Outline

- Part 1: Data definition language
- Part 2: Data manipulation language
 - Queries
 - Insert, delete and update

Relational Query Languages

- A major strength of the relational model: supports simple, powerful querying of data
- Two sublanguages:
 - DDL – Data Definition Language
 - Define and modify schema (at all 3 levels)
 - DML – Data Manipulation Language
 - Queries can be written intuitively
- The DBMS is responsible for efficient evaluation.
 - The key: precise semantics for relational queries
 - Allows the optimizer to extensively re-order operations, and still ensure that the answer does not change
 - Internal cost model drives use of indexes and choice of access paths and physical operators

The SQL Query Language

- The most widely used relational query language
- Originally IBM, then ANSI in 1986
- Current standard is SQL-2008
 - 2003 was last major update: XML, window functions, sequences, auto-generated IDs
 - Not fully supported yet
- SQL-1999 Introduced “Object-Relational” concepts
 - Also not fully supported yet
- SQL92 is a basic subset
 - Most systems support at least this
- PostgreSQL has some “unique” aspects (as do most systems)
- SQL is not synonymous with Microsoft’s “SQL Server”

Part 1: SQL Data Definition (DDL)

- Objects
 - Table
- Commands
 - CREATE
 - ALTER
 - DROP

Tables in SQL

- Created by CREATE TABLE statement

```
CREATE TABLE EMPLOYEE;
```

- Known as base tables
- Attributes ordered by creation order
- Rows not ordered

Description of the Relational Model

- All of the information stored in a Relational Database is held in relations

No other data structures!

- A relation may be thought of as a table

Student

name	id	exam1	exam2
Mounia	891023	12	58
Jane	891024	66	90
Thomas	891025	50	65

- A relation has:
 - a **name**
 - an unchanging set of **columns**; named and typed
 - a time varying set of **rows**

DDL – Create Table

```
CREATE TABLE table_name  
( { column_name data_type  
[ DEFAULT default_expr ] [  
column_constraint [, ... ] ] |  
table_constraint } [, ... ] )
```

- Data Types (PostgreSQL) include:
 - character(n) – fixed-length character string
 - character varying(n) – variable-length character string
 - smallint, integer, bigint, numeric, real, double precision
 - date, time, timestamp, ...
 - serial - unique ID for indexing and cross reference

Data Types

- Numeric

- Integer: INT
- Real: FLOAT

- Character-string

- Fixed length: CHAR(n)
- Varying length: VARCHAR(n)

- DATE

- Has main components YEAR, MONTH, DAY
- Also stores century, hour, minute, second
- Has format DD-MON-YYYY
e.g. 05-FEB-2001

Constraints

- Recall that the schema defines the legal instances of the relations
- Data types are a way to limit the kind of data that can be stored in a table, but they are often insufficient
 - e.g. prices must be positive values
 - uniqueness, etc.
- Can specify constraints on individual columns or on tables

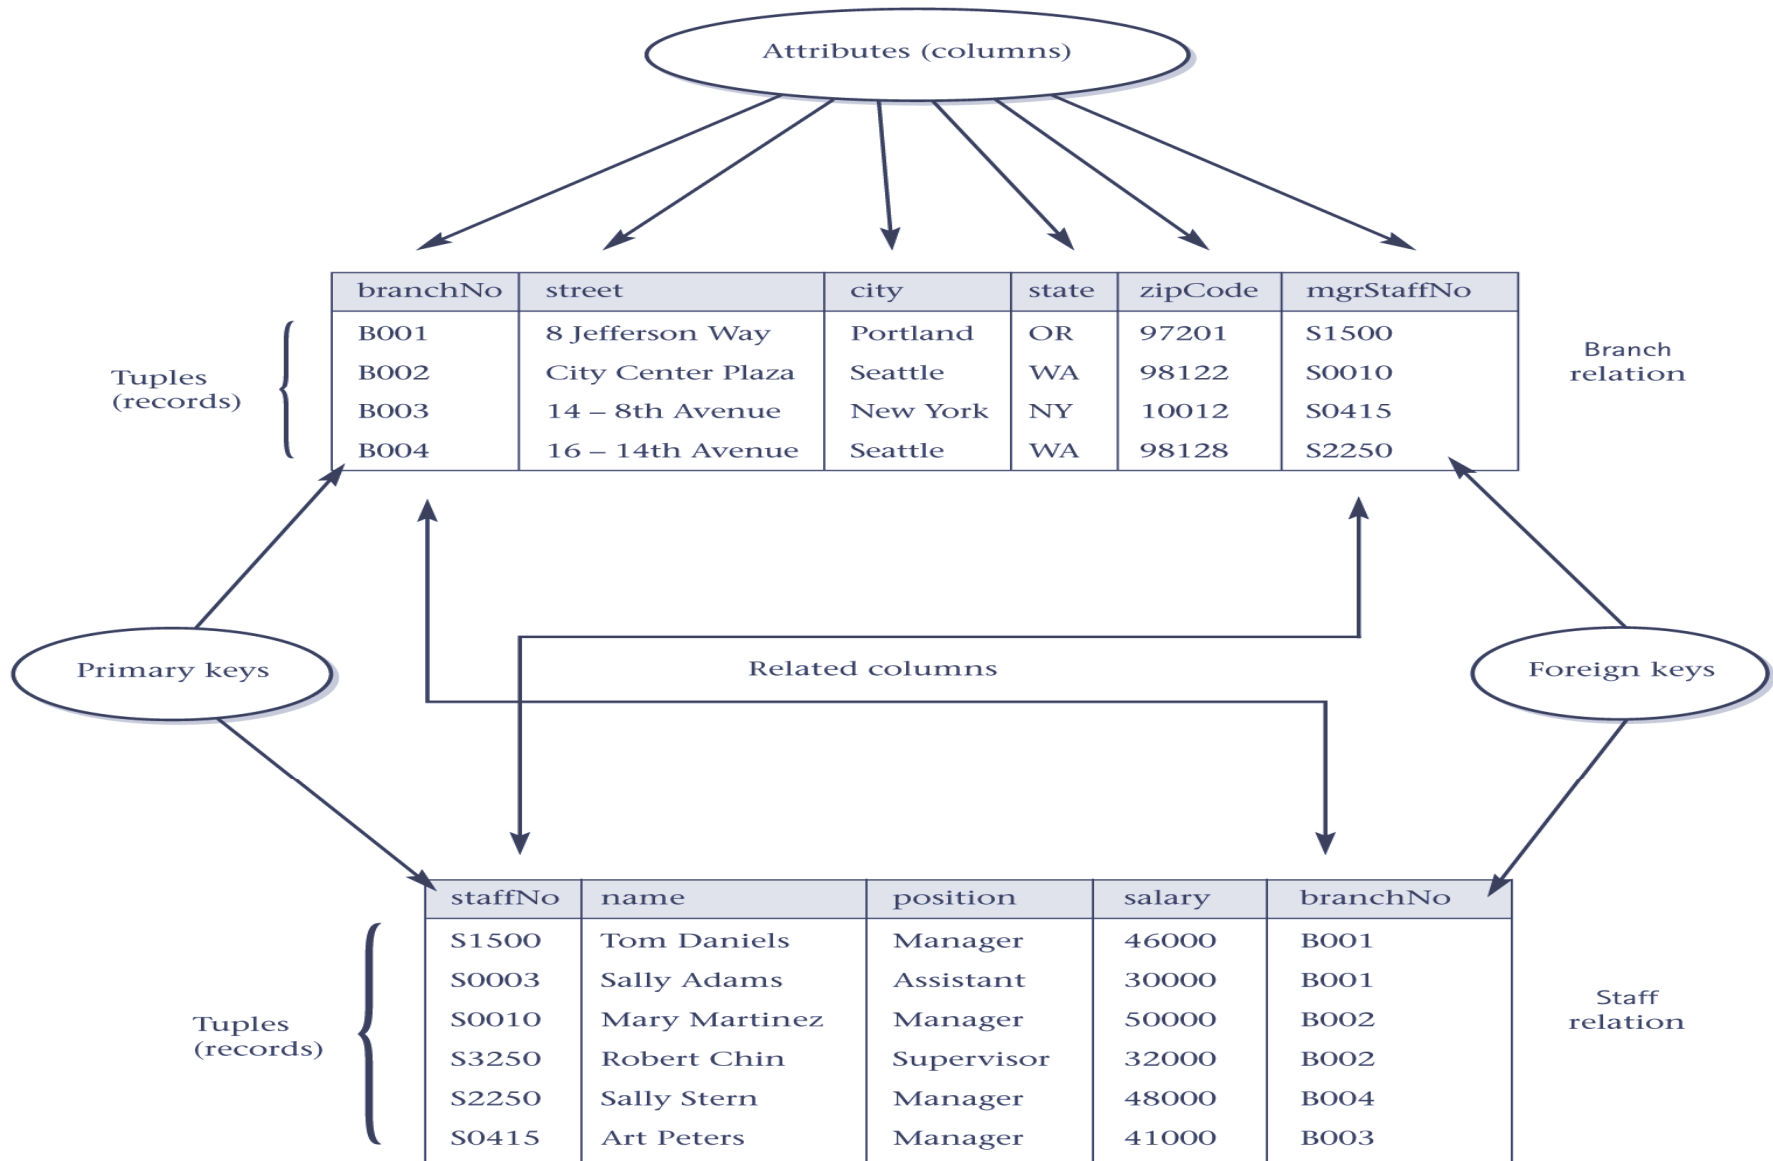
Column (attribute) Constraints

```
[ CONSTRAINT constraint_name ]  
{ NOT NULL | NULL | UNIQUE | PRIMARY KEY |  
  CHECK (expression) |  
  REFERENCES reftable [ ( refcolumn ) ]  
[ ON DELETE action ] [ ON UPDATE action ] }
```

primary key = unique + not null; also used as default target for references (can have at most 1)

references is for foreign keys; *action* is one of:

NO ACTION, CASCADE, SET NULL, SET DEFAULT



Column Constraints: Default Values

- A type of column constraint
- Specified by DEFAULT <value>
- Used if no explicit value assigned to attribute
- NULL unless otherwise stated

Example of Column Constraints

```
CREATE TABLE sp  
(sno VARCHAR(5) NOT NULL REFERENCES s,  
pno VARCHAR(5) NOT NULL REFERENCES p,  
qty INT CHECK (qty >= 0),  
PRIMARY KEY (sno, pno))
```

Table Constraints

- PRIMARY KEY
- UNIQUE (secondary key)
- FOREIGN KEY (referential integrity)
 - Referential integrity constraints can be violated by
 - Insertion or deletion of tuples
 - Foreign key value modified
 - Referential triggered action
 - Can be added to foreign key constraint to cause automatic update: ON UPDATE and ON DELETE
 - Options are SET NULL, CASCADE and SET DEFAULT

Table Constraints

```
CREATE TABLE table_name ( {  
column_name data_type [ DEFAULT default_expr ] [ column_constraint [, ... ] ] |  
table_constraint } [, ... ] )
```

- Table Constraints:

```
[ CONSTRAINT constraint_name ]  
{ UNIQUE ( column_name [, ... ] ) |  
PRIMARY KEY ( column_name [, ... ] ) |  
CHECK ( expression ) |  
FOREIGN KEY ( column_name [, ... ] ) REFERENCES reftable [ ( refcolumn [, ... ] ) ] [  
ON DELETE action ] [ ON UPDATE action ] }
```


Create Table (Examples)

```
CREATE TABLE films (
```

```
    code      CHAR(5) PRIMARY KEY,
```

```
    title     VARCHAR(40),
```

```
    did       DECIMAL(3),
```

```
    date_prod DATE,
```

```
    kind      VARCHAR(10),
```

```
    CONSTRAINT production UNIQUE(date_prod)
```

```
    FOREIGN KEY did REFERENCES distributors
```

```
    ON DELETE NO ACTION
```

```
);
```

```
CREATE TABLE distributors (
```

```
    did       DECIMAL(3) PRIMARY KEY,
```

```
    name      VARCHAR(40)
```

```
    CONSTRAINT con1 CHECK (did > 100 AND name <> ' ')
```

```
);
```



Other DDL Statements

- ALTER TABLE

- use to add/remove columns, constraints, rename things ...

- DROP TABLE

- Compare to “Delete * From Table”

- CREATE/DROP VIEW

- CREATE/DROP INDEX

- GRANT/REVOKE PRIVILEGES

- SQL has an authorization model for saying who can read/modify/delete etc. data and who can grant and revoke privileges!

DROP TABLE

- Option

- CASCADE CONSTRAINT
- e.g. **DROP TABLE** PropertyForRent **CASCADE;**
 - The DROP operation drops all dependent on these objects (and objects dependent on these objects)
- RESTRICT CONSTRAINT
- e.g. **DROP TABLE** PropertyForRent **RESTRICT;**

ALTER TABLE [1]

- Command which allows
 - Adding column

```
ALTER TABLE EMPLOYEE ADD JOB VARCHAR(12);
```

- Dropping column

```
ALTER TABLE EMPLOYEE DROP ADDRESS CASCADE;
```



ALTER TABLE [2]

- Changing column definition (add/drop default)

```
ALTER TABLE DEPARTMENT ALTER MGRSSN  
DROP DEFAULT;
```

```
ALTER TABLE DEPARTMENT ALTER MGRSSN  
SET DEFAULT '11111111';
```



ALTER TABLE [3]

- Adding / dropping table constraints

```
ALTER TABLE EMPLOYEE DROP CONSTRAINT EMPSUPERFK CASCADE;
```

```
ALTER TABLE EMPLOYEE ADD CONSTRAINT EMPSUPERFK;
```

```
FOREIGN KEY(SUPERSSN) REFERENCES EMPLOYEE(SSN)  
ON DELETE SET NULL;
```

Part 2: SQL Data Manipulation (DML)

- Single-table queries are straightforward
 - Example Query 0

To find all 18 year old students, we can write:

```
SELECT *  
FROM Students S  
WHERE S. age=18
```

To find just names and logins, replace the first line:

```
SELECT S. name, S. l o g i n
```



Querying Multiple Relations

- Can specify a join over two tables as follows:

```
SELECT S.name, E.cid  
FROM Students S, Enrolled E  
WHERE S.sid=E.sid AND E.grade='B'
```

sid	cid	grade
53831	Carnatic101	C
53831	Reggae203	B
53650	Topology112	A
53666	History105	B

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@ee	18	3.2

result =

S.name	E.cid
Jones	History105

Note: obviously no
referential integrity
constraints have been
used here

Basic SQL Query

SELECT	[DISTINCT]	<i>target-list</i>
FROM		<i>relation-list</i>
WHERE		<i>qualification</i>

relation-list : A list of relation names

- possibly with a *range-variable* after each name

target-list : A list of attributes of tables in *relation-list*

qualification : Comparisons combined using AND, OR and NOT

- Comparisons are Attr *op* const or Attr1 *op* Attr2, where *op* is one of $= \neq < > \leq \geq$

DISTINCT: optional keyword indicating that the answer should not contain duplicates

- In SQL SELECT, the default is that duplicates are not eliminated! (Result is called a “multiset”)

Query Semantics

- Semantics of an SQL query are defined in terms of the following conceptual evaluation strategy:
 1. do FROM clause: Compute cross-product of tables (e.g. Students and Enrolled)
 2. do WHERE clause: Check conditions, discard tuples that fail, i.e. “selection”
 3. do SELECT clause: Delete unwanted fields, i.e. “projection”
 4. if DISTINCT specified, eliminate duplicate rows

Probably the least efficient way to compute a query!

- An optimizer will find more efficient strategies to get the *same answer*







Step 1 – Cross Product

S.sid	S.name	S.login	S.age	S.gpa	E.sid	E.cid	E.grade
53666	Jones	jones@cs	18	3.4	53831	Carnatic101	C
53666	Jones	jones@cs	18	3.4	53832	Reggae203	B
53666	Jones	jones@cs	18	3.4	53650	Topology112	A
53666	Jones	jones@cs	18	3.4	53666	History105	B
53688	Smith	smith@ee	18	3.2	53831	Carnatic101	C
53688	Smith	smith@ee	18	3.2	53831	Reggae203	B
53688	Smith	smith@ee	18	3.2	53650	Topology112	A
53688	Smith	smith@ee	18	3.2	53666	History105	B

sid	cid	grade
53831	Carnatic101	C
53831	Reggae203	B
53650	Topology112	A
53666	History105	B

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@ee	18	3.2

Step 2 - Discard Tuples that Fail Predicate

S.sid	S.name	S.login	S.age	S.gpa	E.sid	E.cid	E.grade
53666	Jones	jones@acs	18	3.4	53831	Camatic101	C
53666	Jones	jones@acs	18	3.4	53832	Reggae203	
53666	Jones	jones@acs	18	3.4	53650	Topology112	A
	Jones	jones@acs	18	3.4		History105	
53688	Smith	smith@ee	18	3.2	53831	Camatic101	C
53688	Smith	smith@ee	18	3.2	53831	Reggae203	
53688	Smith	smith@ee	18	3.2	53650	Topology112	A
53688	Smith	smith@ee	18	3.2	53666	History105	

```
SELECT S. name, E. ci d
FROM Students S, Enrolled E
WHERE S. si d=E. si d AND E. grade= 'B'
```

Step 3 - Discard Unwanted Columns

S.sid	S.name	S.login	S.age	S.gpa	E.sid	E.cid	E.grade
53666	Jones	jones@cs	18	3.4	53831	Carnatic101	C
53666	Jones	jones@cs	18	3.4	53832	Reggae203	B
53666	Jones	jones@cs	18	3.4	53650	Topology112	A
53666	Jones	jones@cs	18	3.4	53666	History105	B
53688	Smith	smith@ee	18	3.2	53831	Carnatic101	C
53688	Smith	smith@ee	18	3.2	53831	Reggae203	B
53688	Smith	smith@ee	18	3.2	53650	Topology112	A
53688	Smith	smith@ee	18	3.2	53666	History105	B

```
SELECT S.name, E.cid  
FROM Students S, Enrolled E  
WHERE S.sid=E.sid AND E.grade='B'
```

Now the Details...

We will use these relations in our examples

Reserves

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
95	103	11/12/96

Sailors

<u>sid</u>	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
95	Bob	3	63.5

Boats

<u>bid</u>	bname	color	
101	Interlake	blue	
102	Interlake	red	
103	Clipper	green	
104	Marine	red	

Example Schemas (in SQL DDL)

```
CREATE TABLE Sailors (sid INTEGER, sname CHAR(20), rating INTEGER,  
age REAL, PRIMARY KEY sid)
```

```
CREATE TABLE Boats (bid INTEGER, bname CHAR (20), color CHAR(10),  
PRIMARY KEY bid)
```

```
CREATE TABLE Reserves (sid INTEGER, bid INTEGER, day DATE,  
PRIMARY KEY (sid, bid, day),  
FOREIGN KEY sid REFERENCES Sailors,  
FOREIGN KEY bid REFERENCES Boats)
```

Another Join Query

```
SELECT  sname
FROM    Sailors, Reserves
WHERE   Sailors.sid=Reserves.sid
AND     bid=103
```

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
95	Bob	3	63.5	22	101	10/10/96
95	Bob	3	63.5	95	103	11/12/96

Some Notes on Aliasing

- Can associate an alias with the tables in the FROM clause
 - saves writing, makes queries easier to understand
- Needed when ambiguity could arise
 - for example, if same table used multiple times in same FROM (called a “self-join”)

```
SELECT sname
FROM Sailors, Reserves
WHERE Sailors.sid=Reserves.sid AND bid=103
```

Can be
rewritten using
range variables as:

```
SELECT S.sname
FROM Sailors S, Reserves R
WHERE S.sid=R.sid AND bid=103
```

More Notes

- Here's an example where aliases are required (self-join example):

```
SELECT  x. sname,  x. age,  y. sname,  y. age
FROM    Sai l ors x, Sai l ors y
WHERE   x. age > y. age
```

Note that target list can be replaced by “*” if you don't want to do a projection:

```
SELECT  *
FROM    Sai l ors x
WHERE   x. age > 20
```

Find sailors who've reserved at least one boat

```
SELECT  S. sid  
FROM    Sailors S, Reserves R  
WHERE   S. sid=R. sid
```

Would adding DISTINCT to this query make a difference?

What is the effect of replacing *S.sid* by *S.sname* in the SELECT clause?

- Would adding DISTINCT to this variant of the query make a difference?

Expressions

- Can use arithmetic expressions in SELECT clause (plus other operations we will discuss later)

Use **AS** to provide column names (Example - q1b)

```
SELECT S. age, S. age-5 AS age1, 2*S. age AS age2
FROM   Sai l o r s S
WHERE  S. sname = 'dust i n'
```

Can also have expressions in WHERE clause:

```
SELECT  S1. sname AS name1, S2. sname AS name2
FROM    Sai l o r s S1, Sai l o r s S2
WHERE   2*S1. rati ng = S2. rati ng - 1
```

No WHERE Clause

- No condition on tuple selection
- Example - query 9
- More than one relation in FROM clause means cross product
- Example - query 10
- Similar to relational algebra cross product - PROJECT combination

Use of Asterisk

- Used to retrieve all attribute values in SELECT clause
- Examples - queries 1C, 1D, 10A

String Operations

- SQL also supports some string operations
- Example - queries 12, 12A
- “LIKE” is used for string matching

```
SELECT  S. age, age1=S. age-5, 2*S. age AS age2  
FROM    Sailors S  
WHERE   S. sname LIKE 'B_%B'
```

‘_’ stands for any one character and ‘%’ stands for 0 or more arbitrary characters.

*Find 'sid's of sailors who have reserved a **red** or a **green** boat (Example - queries 11, 11A)*

- **UNION** can be used to compute the union of any two **union-compatible** sets of tuples (which are themselves the result of SQL queries)

Vs.

```
SELECT DISTINCT R. si d
FROM Boats B, Reserves R
WHERE R. bi d=B. bi d AND
(B. col or= 'red' OR B. col or= 'green' )
```

(note:
UNION
eliminates
duplicates
by default.
Override w/
UNION ALL)

```
SELECT R. si d
FROM Boats B, Reserves R
WHERE R. bi d=B. bi d AND B. col or= 'red'  UNI ON
SELECT R. si d
      FROM Boats B, Reserves R
      WHERE R. bi d=B. bi d AND
            B. col or= 'green'
```


Nested Queries

- Powerful feature of SQL: WHERE clause can itself contain an SQL query! (Example - query 16)

- Actually, so can FROM and HAVING clauses

Names of sailors who've reserved boat #103:

```
SELECT  S. sname
FROM    Sailors S
WHERE   S. sid IN (SELECT R. sid
                  FROM   Reserves R
                  WHERE  R. bid=103)
```

To find sailors who've **not** reserved #103, use NOT IN

To understand semantics of nested queries:

- **think** of a nested loops evaluation: *For each Sailors tuple, check the qualification by computing the subquery*

Nested Queries with Correlation

Find names of sailors who've reserved boat #103:

```
SELECT  S. sname
FROM    Sailors S
WHERE   EXISTS (SELECT  *
                  FROM    Reserves R
                  WHERE   R. bi d=103  AND  S. si d=R. si d)
```



- EXISTS is another set comparison operator, like IN
- Can also specify NOT EXISTS
- Subquery must be recomputed for each Sailors tuple
 - Think of subquery as a function call that runs a query!

Other Comparison Operators

- $>$, $>=$, $<$, $<=$, $<>$
- Can be used with ANY, SOME, ALL

```
SELECT  LNAME, FNAME
FROM    EMPLOYEE
WHERE   SALARY > ALL      (SELECT SALARY
                           FROM    EMPLOYEE
                           WHERE   DNO = 5);
```

```
SELECT LNAME, SALARY
FROM EMPLOYESS
WHERE SALARY > ALL (1600, 2999);
```

Arithmetic and Other Operators

- Standard arithmetic operators can be applied
- Example - query 13
- String concatenation ||
- Numeric value range BETWEEN
- Example - query 14
- Ordering by value of one or more attributes
- Example - query 15

Aggregate Operators

- Significant extension of relational algebra
- Example queries: 19, 20, 21, 22, 23, 5

```
SELECT COUNT(*)  
FROM Sailors S
```

```
SELECT AVG(S. age)  
FROM Sailors S  
WHERE S. rating=10
```

```
SELECT COUNT(DISTINCT S. rating)  
FROM Sailors S  
WHERE S. sname= 'Bob'
```

COUNT (*)
COUNT ([DISTINCT] A)
SUM ([DISTINCT] A)
AVG ([DISTINCT] A)
MAX (A)
MIN (A)

single column

Find name and age of the oldest sailor(s)

- The first query is incorrect!
 - Returns the sname of each sailor along with the (same) age of the oldest sailor

```
SELECT S. sname, MAX (S. age)  
FROM Sai l ors S
```

- The second query is correct!

```
SELECT S. sname, S. age  
FROM   Sai l ors S  
WHERE  S. age =  
        (SELECT MAX(S2. age)  
         FROM Sai l ors S2)
```

GROUP BY and HAVING

- So far, we have applied aggregate operators to all (qualifying) tuples
 - Sometimes, we want to apply them to each of several *groups* of tuples
- Consider: Find the age of the youngest sailor for each rating level
 - In general, we don't know how many rating levels exist, and what the rating values for these levels are!
 - Suppose we know that rating values go from 1 to 10; we can write 10 queries that look like this (!):

```
For  $i = 1, 2, \dots, 10$ :  
SELECT  MIN (S. age)  
FROM    Sailors S  
WHERE   S. rating = i
```

Queries with GROUP BY

- To generate values for a column based on groups of rows, use **aggregate** functions in SELECT statements with the GROUP BY clause

```
SELECT  [DI STI NCT]  target-list
FROM    relation-list
[WHERE  qual i f i c a t i o n]
GROUP BY grouping-list
```

The *target-list* contains:

- I. list of column names &
- II. terms with aggregate operations e.g. MIN (S.age)

column name list (I) can contain only attributes from the *grouping-list*

GROUP BY Examples (Example - queries 24&25)

For each rating, find the average age of the sailors

```
SELECT S.rating, AVG (S.age)
FROM Sailors S
GROUP BY S.rating
```

For each rating find the age of the youngest
sailor with age \geq 18

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
WHERE S.age >= 18
GROUP BY S.rating
```

Conceptual Evaluation

- The cross-product of **relation-list** is computed, tuples that fail **qualification** are discarded, 'unnecessary' fields are deleted, and the remaining tuples are partitioned into groups by the value of attributes in **grouping-list**
- One answer tuple is generated per qualifying group

```

SELECT  S.rating,  MIN (S.age)
FROM    Sailors S
WHERE   S.age >= 18
GROUP BY  S.rating

```

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

1. Form cross product

rating	age
1	33.0
7	45.0
7	35.0
8	55.5
10	35.0

2. Delete unneeded columns,
rows; form groups

3. Perform Aggregation

rating	age
1	33.0
7	35.0
8	55.0
10	35.0

Queries with GROUP BY and HAVING

SELECT	[DISTINCT] <i>target-list</i>
FROM	<i>relation-list</i>
WHERE	<i>qualification</i>
GROUP BY	<i>grouping-list</i>
HAVING	<i>group-qualification</i>

- Use the HAVING clause with the GROUP BY clause to restrict which group-rows are returned in the result set

(Example - queries 26, 28)

*Find the age of the youngest sailor with age ≥ 18 ,
for each rating with at least 2 such sailors*

```
SELECT  S.rating, MIN(S.age)
FROM    Sailors S
WHERE   S.age >= 18
GROUP BY S.rating
HAVING COUNT (*) > 1
```

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

rating	age
1	33.0
7	45.0
7	35.0
8	55.5
10	35.0

rating	m-age	count
1	33.0	1
7	35.0	2
8	55.0	1
10	35.0	1

rating	
7	35.0

Answer

Summary of SQL Queries

SELECT <attribute and function list>

FROM <table list>

[WHERE <condition>]

[GROUP BY <grouping attribute(s)>]

[HAVING <group condition>]

[ORDER BY <attribute list>]

INSERT

```
INSERT [INTO] table_name [(column_list)]  
VALUES (value_list)
```

```
INSERT [INTO] table_name [(column_list)]  
<select statement>
```

```
INSERT INTO Boats VALUES ( 105, 'Clipper' , 'purple' )
```

```
INSERT INTO Boats (bid, color) VALUES (99, 'yellow' )
```

- You can also do a “bulk insert” of values from one table into another (must be type compatible):

```
INSERT INTO TEMP(bid)
```

```
SELECT r.bid FROM Reserves R WHERE r.sid = 22
```

(Example - update 1)

DELETE & UPDATE

DELETE [FROM] <i>table_name</i> [WHERE <i>qualification</i>]
--

DELETE FROM Boats WHERE color = 'red'

DELETE FROM Boats b

WHERE b. bid =

(SELECT r.bid FROM Reserves R WHERE r.sid = 22)

- Can also modify tuples using UPDATE statement:

UPDATE Boats

SET Color = "green"

WHERE bid = 103;

(Example - updates 6, 5)

Views

```
CREATE VIEW view_name  
AS select_statement
```

- Makes development simpler
- Often used for security
- Not instantiated - makes updates tricky

```
CREATE VIEW Reds  
AS SELECT B.bid, COUNT (*) AS scount  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid AND B.color= 'red'  
GROUP BY B.bid
```

```
CREATE VIEW Reds AS
  SELECT B.bid, COUNT (*) AS scount
  FROM Boats B, Reserves R
  WHERE R.bid=B.bid AND B.color= 'red'
  GROUP BY B.bid
```

b.bid	scount	<i>Reds</i>
102	1	

Querying and Deleting Views

- SQL queries can be specified on views

(Example - query view 1)

- Deleted using DROP VIEW

(Example - view 1A)

Assertions [1]

- More general constraints can be specified via declarative assertions
- Created using CREATE ASSERTION
- Deleted using DROP ASSERTION

Assertions [2]

```
CREATE ASSERTION SALARY_CONSTRAINT
CHECK (NOT EXISTS (SELECT * FROM EMPLOYEE E,
                    EMPLOYEE M, DEPARTMENT D
                    WHERE E.SALARY > M.SALARY AND
                          E.DNO = D.DNUMBER AND
                          D.MGRSSN = M.SSN) );
```

- CHECK clause can be used with CREATE DOMAIN statement:

```
CREATE DOMAIN D_NUM AS INTEGER
CHECK (D_NUM > 0 AND D_NUM < 21);
```

```

SELECT r.sid, b.bid, b.name
FROM Reserves r FULL OUTER JOIN Boats b
ON r.bid = b.bid

```

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
95	103	11/12/96

<u>bid</u>	bname	color	
101	Interlake	blue	
102	Interlake	red	
103	Clipper	green	
104	Marine	red	

r.sid	b.bid	b.name
22	101	Interlake
	102	Interlake
95	103	Clipper
	104	Marine

Note: in this case it is the same as the ROJ because bid is a foreign key in reserves, so all reservations must have a corresponding tuple in boats

Division in SQL

Find names of sailors who've reserved all boats

Example in book, not using EXCEPT:

```
SELECT S.sname    Sailors S such that ...
FROM   Sailors S
WHERE  NOT EXISTS (SELECT B.bid
there is no boat B   FROM Boats B
                     WHERE NOT EXISTS (SELECT R.bid
that doesn't have ...   FROM Reserves R
                     WHERE R.bid=B.bid
a Reserves tuple showing S reserved B   AND R.sid=S.sid))
```

*Find the number of reservations for each **red** boat*

- Grouping over a join of two relations

```
SELECT B.bid, COUNT(*) AS scout  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid  
      AND B.color= 'red'  
GROUP BY B.bid
```



```

SELECT B.bid, COUNT (*) AS scount
FROM Boats B, Reserves R
WHERE R.bid=B.bid AND B.color= 'red'
GROUP BY B.bid

```

b.bid	b.color	r.bid
101	blue	101
102	red	101
103	green	101
104	red	101
101	blue	102
102	red	102
103	green	102
104	red	102

1

b.bid	b.color	r.bid
102	red	102

2

b.bid	scount
102	1

answer

Null Values – 3 Valued Logic

(null > 0) is null

(null + 1) is null

(null = 0) is null

null AND true is null

AND	T	F	Null
T	T	F	Null
F	F	F	F
NULL	Null	F	Null

OR	T	F	Null
T	T	T	T
F	T	F	Null
NULL	T	Null	Null

Joins

```
SELECT (column_list)
FROM   table_name
       [INNER | {LEFT|RIGHT|FULL} OUTER] JOIN table_name
       ON qualification_list
WHERE ...
```

- Explicit join semantics needed unless it is an INNER join (INNER is default)

Inner Join

- Only the rows that match the search conditions are returned

```
SELECT s.sid, s.name, r.bid  
FROM Sailors s INNER JOIN Reserves r  
ON s.sid = r.sid
```

Returns only those sailors who have reserved boats

SQL-92 also allows:

```
SELECT s.sid, s.name, r.bid  
FROM Sailors s NATURAL JOIN Reserves r
```

- “NATURAL” in SQL means equi-join for each pair of attributes with the same name (may need to rename with “AS”)

```
SELECT s.sid, s.name, r.bid
FROM Sailors s INNER JOIN Reserves r
ON s.sid = r.sid
```

<u>sid</u>	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
95	Bob	3	63.5

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
95	103	11/12/96

s.sid	s.name	r.bid
22	Dustin	101
95	Bob	103

Left Outer Join

- Left Outer Join returns all matched rows, plus all unmatched rows from the table on the left of the join clause (use nulls in fields of non-matching tuples)

```
SELECT s.sid, s.name, r.bid  
FROM Sailors s LEFT OUTER JOIN Reserves r  
ON s.sid = r.sid
```

Returns all sailors & information on whether they have reserved boats

```

SELECT s.sid, s.name, r.bid
FROM Sailors s LEFT OUTER JOIN Reserves r
ON s.sid = r.sid

```

<u>sid</u>	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
95	Bob	3	63.5

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
95	103	11/12/96

s.sid	s.name	r.bid
22	Dustin	101
95	Bob	103
31	Lubber	

Right Outer Join

- Right Outer Join returns all matched rows, plus all unmatched rows from the table on the right of the join clause

```
SELECT r.sid, b.bid, b.name  
FROM Reserves r RIGHT OUTER JOIN Boats b  
ON r.bid = b.bid
```

Returns all boats & information on which ones are reserved


```

SELECT r.sid, b.bid, b.name
FROM Reserves r RIGHT OUTER JOIN Boats b
ON r.bid = b.bid

```

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
95	103	11/12/96

<u>bid</u>	bname	color	
101	Interlake	blue	
102	Interlake	red	
103	Clipper	green	
104	Marine	red	

r.sid	b.bid	b.name
22	101	Interlake
	102	Interlake
95	103	Clipper
	104	Marine

Full Outer Join

- Full Outer Join returns all (matched or unmatched) rows from the tables on both sides of the join clause

```
SELECT r.sid, b.bid, b.name  
FROM Reserves r FULL OUTER JOIN Boats b  
ON r.bid = b.bid
```

Returns all boats & all information on reservations

Find names of sailors who've reserved all boats

Example in book, not using EXCEPT:

```
SELECT S.sname
FROM Sailors S    Sailors S such that ...
WHERE NOT EXISTS (SELECT B.bid
                  FROM Boats B    there is no boat B without ...
                  WHERE NOT EXISTS (SELECT R.bid
                                   a Reserves tuple showing S reserved B
                                   FROM Reserves R
                                   WHERE R.bid=B.bid
                                   AND R.sid=S.sid))
```

Find names of sailors who've reserved all boats

Can you do this using Group By and Having?

```
SELECT S.sname
FROM Sailors S, reserves R
WHERE S.sid = R.sid
GROUP BY S.sname, S.sid
HAVING
    COUNT(DISTINCT R.bid) =
        (Select COUNT (*) FROM Boats)
```

Note: must have both sid and name in the GROUP BY clause. Why?

Sailors who have reserved all boats

```
SELECT S.name
FROM Sailors S, reserves R
WHERE S.sid = R.sid
GROUP BY S.name, S.sid
HAVING COUNT(DISTINCT R.bid) =
  (Select COUNT (*) FROM Boats)
```

Sailors

sid	sname	rating	age
1	Frodo	7	22
2	Bilbo	2	39
3	Sam	8	27

Boats

bid	bname	color
101	Nina	red
102	Pinta	blue
103	Santa Maria	red

Reserves

sid	bid	day
1	102	9/12
2	102	9/12
2	101	9/14
1	102	9/10
2	103	9/13

sname	sid	bid
Frodo	1	102
Bilbo	2	101
Bilbo	2	102
Frodo	1	102
Bilbo	2	103

sname	sid	count
Frodo	1	1
Bilbo	2	3

count
3

sname	sid	bid
Frodo	1	102,102
Bilbo	2	101, 102, 103