

Overview

- How to deal with situations where multiple users access and modify entries in the database at the same time?
- Problems that can occur (conflicts)
 - How to formally represent multiple database accesses?
 - What are the typical problems?
- Ways of addressing conflicts
 - Locking
 - Serialisation

Transactions

- An action or a series of actions carried out by a single user or application program, which read/update the contents of the DB
- A transaction is a logical unit of DB processing, consisting of one or more DB access operations
- Transaction boundaries may be specified implicitly or explicitly
- Transactions are recorded in system log, kept on disk

Example Transactions

Successful

```
begin_transaction;  
read_item(X);  
X:=X-10;  
write_item(X);read_item(Y);  
Y:=Y+5;  
write_item(Y);  
end_transaction;  
commit;
```

Unsuccessful

```
begin_transaction;  
read_item(X);  
X:=X-10;  
write_item(X);read_item(Y);  
[transaction fails]  
abort;
```

Concurrency

- A number of operations that are executed while overlapping in time
- Concurrency is particularly important if the different operations use the same data
- Conflicts can arise
- A number of solutions
 - Serialisation
 - Locking

Transaction Properties [1]

- Must hold for every transaction for the DB to remain stable
- **ACID** properties
 - **A**tomicity
 - A transaction should be treated as an indivisible unit
 - Managed by transaction recovery subsystem
 - **C**onsistency preservation
 - A transaction must transform the database from one consistent state to another consistent state - only valid data will be written to the DB
 - Managed by programmers / DBMS module

Transaction Properties [2]

- **I**solation
 - Transactions should execute independently of each other
 - Managed by concurrency control subsystem
- **D**urability
 - Effects of a successful transaction must be permanently recorded in the DB
 - Managed by recovery subsystem

Concurrency Control

- Concurrency control is necessary because of
 - Lost update problem
 - Uncommitted dependency (or dirty read) problem
 - Inconsistent analysis problem

Lost Update Problem

- An apparently completed update by one user can be overridden by another user
- While T1 (transaction 1) reads the value of an item, the value of that item is changed by T2 (transaction 2)
- This can lead to situations where one of the changes (updates) of one transaction are disregarded (lost)

EXAMPLE

Time	T1	T2	X
t ₁		begin_transaction	100
t ₂	begin_transaction	read(X)	100
t ₃	read(X)	X:=X+100	100
t ₄	X:=X-10	write(X)	200
t ₅	write(X)	commit	90
t ₆	commit		90

Uncommitted Dependency Problem

- The uncommitted dependency problem occurs when T3 is allowed to use the result of another transaction (T4) before T4 has committed the changes
- The reason for not committing vary (cancelled by the user, connection problems, system crashes, etc.)
- Failure to commit causes a rollback, but other transactions are unaware of the rollback

EXAMPLE

Time	T3	T4	X
t ₁		begin_transaction	100
t ₂		read(X)	100
t ₃		X:=X+100	100
t ₄	begin_transaction	write(X)	200
t ₅	read(X)	...	200
t ₆	X:=X-10	rollback	100
t ₇	write(X)		190
t ₈	commit		190

Inconsistent Analysis Problem

- If one transaction (T6) reads in several items from the database while some of the values are updated by another transaction (T5)
- This leads to a situation where some of the values T6 uses are not (yet) updated by T5 but some of the items have been updated

EXAMPLE

Time	T5	T6	X	Y	Z	sum
t ₁		begin_transaction	100	50	25	
t ₂	begin_transaction	sum=0	100	50	25	0
t ₃	read(X)	read(X)	100	50	25	0
t ₄	X:=X-10	sum:=sum+X	100	50	25	100
t ₅	write(X)	read(Y)	90	50	25	100
t ₆	read(Z)	sum:=sum+Y	90	50	25	150
t ₇	Z:=Z+10		90	50	25	150
t ₈	write(Z)		90	50	35	150
t ₉	commit	read(Z)	90	50	35	150
t ₁₀		sum:=sum+Z	90	50	35	185
t ₁₁		commit	90	50	35	185

Schedules and Serialisability

- The objective of a concurrency control is to schedule transactions in such a way as to avoid any interference between them
- One solution is to allow only one transaction at a time
- Other solution is for scheduled transactions to work concurrently. How do we order them in time?

Schedules

- A sequence of operations by a set of concurrent transactions that preserves the order of the operations in each individual transaction
- Two operations conflict if
 - They belong to different transactions **AND**
 - They access the same item (e.g. X, Y) **AND**
 - At least one of the operations is a `write_item(X)`

Schedule Criteria [1]

- Complete
 - All operations from each of the transactions is present (read, write, etc.)
 - Commit or abort operation must be last in each transaction
 - Any pair of operations from same transaction appear in correct order
 - For any pair of conflicting operations, one must occur before the other in the schedule
- Partial order of operations
 - Two non-conflicting operations may occur simultaneously

Schedule Criteria [2]

- Recoverability
 - Should not be necessary to rollback, i.e. undo write operations to DB after commit point
 - Transaction only commits when all other transactions that are writing to common data item have committed
- Avoidance of cascading rollback
 - All transactions only read items written by committed transactions
- Strictness
 - Transactions cannot read or write items until last transaction to write item has committed or aborted

Example of Nonrecoverable Schedule

Time	T1	T2
t ₁	begin_transaction	
t ₂	read(X)	
t ₃	X:=X+10	
t ₄	write(X)	begin_transaction
t ₅		read(X)
t ₆		X:=X*1.1
t ₇		write(X)
t ₈		read(Y)
t ₉		Y:=Y*1.1
t ₁₀		write(Y)
t ₁₁	read(Y)	commit
t ₁₂	Y:=Y-100	
t ₁₃	write(Y)	
t ₁₄	commit	

Types of Schedules

- Serial schedule
 - All operations from each transaction are executed consecutively, without operations from different transactions interleaving
- Non-serial schedule
 - Operations from different transactions are interleaved

Schedule Equivalence

- Result equivalent
 - Produces same final DB state
- View equivalent:
- For two schedules S and S' it is the case that
 - If T in S reads the initial value of X then T in S' reads the initial value of X
 - If T_1 reads the value of X , written to by T_2 then T_1 must also read the value of X written to by T_2 in S'
 - If the last write operation to X in S was done by T_1 , then T_1 also to be the last transaction to write to X in S'

Example of Serial Schedule

Time	T1	T2
t ₁	begin_transaction	
t ₂	read(X)	
t ₃	write(X)	
t ₄	read(Y)	
t ₅	write(Y)	
t ₆	commit	
t ₇		begin_transaction
t ₈		read(X)
t ₉		write(X)
t ₁₀		read(Y)
t ₁₁		write(Y)
t ₁₂		commit

Example of Non-Serial, Equivalent Schedules

Time	T1	T2	T1	T2
t ₁	begin_trans		begin_trans	
t ₂	read(X)		read(X)	
t ₃	write(X)		write(X)	
t ₄		begin_trans		begin_trans
t ₅		read(X)		read(X)
t ₆		write(X)	read(Y)	
t ₇	read(Y)			write(X)
t ₈	write(Y)		write(Y)	
t ₉	commit		commit	
t ₁₀		read(Y)		read(Y)
t ₁₁		write(Y)		write(Y)
t ₁₂		commit		commit

Serialisable Schedules [1]

- A non-serial schedule is called serialisable if
 - there is a way in which the transactions can be executed concurrently without interfering with one another
 - and thereby produce a database state that could be produced by a serial execution
- If a set of transactions executes concurrently, we say that the schedule is correct if it produces the same results as some serial execution

Serialisable Schedules [2]

- In Serialisability, the ordering of read and write operations is important
 - If two transactions only read an item, they do not conflict and the **order is not important**
 - If two transactions either read or write completely separate items, they do not conflict and **the order is not important**
 - If one transaction writes an item and another reads or writes the same item, they could conflict and **the order is important**

Conflict Serialisability


- Two transactions T1 and T2 are in conflict if an operation in T1 and an operation in T2 forces a temporal order between the transactions
- If two transaction can be serialised while maintaining these conflicts (temporal orders), they are called conflict serialisable

Precedence Graphs

- Precedence graphs allow one to visualize conflicts between transactions
- A precedence graph has:
 - A node for each transaction
 - A directed edge $T_i \rightarrow T_j$, if T_j reads the value of an item written to by T_i
 - A directed edge $T_i \rightarrow T_j$, if T_j writes to an item after it has been read by T_i
 - A directed edge $T_i \rightarrow T_j$, if T_j writes to an item after it has been written to by T_i
- If there is a cycle in the graph, the transactions are not conflict serialisable

Example of Precedence Graph

Time	T1	T2	T3
t ₁	begin_trans		
t ₂	read(X)		
t ₃		begin_trans	
t ₄		write(X)	
t ₅		commit	
t ₆	write(X)		
t ₇	commit		
t ₈			begin_trans
t ₉			write(X)
t ₁₀			commit



Concurrency Control Techniques

- Needed to make a schedule
- Locking
- Deadlocks
- Timestamps

Locking [1]

- A lock prevents several transactions from accessing and/or changing items while another transaction accesses and/or changes it
- Lock
 - Variable describing status of an item
 - Information held in a lock table
- Danger of deadlock
 - Each transaction in a set of transactions is waiting for an item locked by another transaction in the same set

Locking [2]

- Types of locks
 - Binary
 - Two possible states: locked or unlocked
 - Two operations: `lock_item(X)` , `unlock_item(X)`
 - Shared /Exclusive
 - Multiple states
 - Three operations:
`read_lock(X)` , `write_lock(X)` , `unlock(X)`

Binary Locks

- Transaction must lock data item before `read_item` or `write_item` operations
- Transaction must unlock data item after finishing with it
- No two transactions can access same data item concurrently

Shared /Exclusive Locks

- Shared lock: If a transaction has a shared lock on an item, it can read the item but not update it
- Exclusive Lock: If a transaction has an exclusive lock on an item, it can both read and update it
- Some systems allow shared locks to be upgraded to exclusive locks
- Similarly, sometimes exclusive locks can be downgraded to shared locks

Locking Schedule

Time	T1	T2
$t_{1,2}$	<code>write_lock(X);read(X)</code>	
t_3	<code>X:=X+100</code>	
$t_{4,5}$	<code>write(X);unlock(X)</code>	<code>begin_transaction</code>
$t_{6,7}$		<code>write_lock(X);read(X)</code>
t_8		<code>X:=X*1.1</code>
$t_{9,10}$		<code>write(X);unlock(X)</code>
$t_{11,12}$		<code>write_lock(Y);read(Y)</code>
t_{13}		<code>Y:=Y*1.1</code>
$t_{14,15}$		<code>write(Y);unlock(Y)</code>
$t_{16,17}$	<code>write_lock(Y);read(Y)</code>	<code>commit</code>
t_{18}	<code>Y:=Y-100</code>	
$t_{19,20}$	<code>write(Y);unlock(Y)</code>	
t_{21}	<code>commit</code>	

If $X=100$ and $Y=400$, if $T_1 < T_2$ then $X=220$, $Y=330$, and if $T_2 < T_1$ then $X=210$, $Y=340$, but the schedule above gives: $X=220$ and $Y=340$

Two-phase Locking (2PL)

- The previous example shows that locking does not always guarantee serialisability
- 2PL does ensure serialisability
- Principle of 2PL
 - Every transaction must lock an item (read or write) before accessing it
 - Once a lock has been released, no new items can be locked
- 2PL has two phases: the growing phase and the shrinking phase
- Growing phase: all locks are acquired but cannot release any
- Shrinking phase: all locks are released but cannot acquire any

Preventing the Lost Update Problem

Time	T1	T2	X
t ₁		begin_transaction	100
t ₂	begin_transaction	write_lock(X)	100
t ₃	write_lock(X)	read(X)	100
t ₄	WAIT	X:=X+100	100
t ₅	WAIT	write(X)	200
t ₆	WAIT	commit/unlock(X)	200
t ₇	read(X)		200
t ₈	X:=X-10		200
t ₉	write(X)		190
t ₁₀	commit/unlock(X)		190

Preventing the Uncommitted Dependency Problem

Time	T3	T4	X
t ₁		begin_transaction	100
t ₃		write_lock(X)	100
t ₄		read(X)	100
t ₅	begin_transaction	X:=X+100	100
t ₄	write_lock(X)	write(X)	200
t ₅	WAIT	rollback/unlock(X)	100
t ₄	read(X)		100
t ₆	X:=X-10		100
t ₇	write(X)		90
t ₈	commit/unlock(X)		90

What are Deadlocks?

- A deadlock occurs if several transaction are waiting for each other to unlock a data item
- General solutions to deadlocks
 - Timeouts
 - Deadlock *prevention*
 - Deadlock *detection* and *recovery*

Example of Deadlocks

Time	T1	T2
t ₁	begin_transaction	
t ₂	write_lock(X)	begin_transaction
t ₃	read(X)	write_lock(Y)
t ₄	X:=X-10	read(Y)
t ₅	write(X)	Y:=Y+100
t ₆	write_lock(Y)	write(Y)
t ₇	WAIT	write_lock(X)
t ₈	WAIT	WAIT
t ₉	WAIT	WAIT
t ₁₀

Deadlock Prevention

- Timeouts
 - A transaction waits a specified amount of time for a lock
 - If the lock is not granted, the transaction is aborted
- Deadlock prevention
 - Order transactions by time, abort younger one, while keeping track of the original time the younger transaction was started
 - Conservative 2PL: All data items have to be locked at the beginning of a transaction
 - Disadvantage of conservative 2PL: locks are held longer; time consuming to determine whether all required locks are free; longer waits for all required locks

Timestamping

- A timestamp is a unique identifier (created by the DBMS) that indicates the relative starting time of transactions
- Timestamping is a concurrency protocol that order transactions, such that transactions with smaller timestamps get priority in the event of a conflict

Summary

- Transactions
- Concurrency problems (lost update problem etc.)
- Schedules (serial, non-serial, serialisability)
- Conflict serialisability and precedence graphs
- Locking (2PL)
- Deadlocks