

ECS524

Application Layer


Prof. Steve Uhlig
steve.uhlig@qmul.ac.uk
Office: Eng.202

Dr. Felix Cuadrado
felix.cuadrado@qmul.ac.uk
Office: Eng 202

Slides

Disclaimer:
Some of the slides' content is
borrowed directly from those
provided by the authors of the
textbook. They are available from

[http://www-
net.cs.umass.edu/kurose-ross-ppt-
6e](http://www-net.cs.umass.edu/kurose-ross-ppt-6e)



*Computer
Networking: A Top
Down Approach*

6th edition


Jim Kurose, Keith Ross
Addison-Wesley
March 2012

The Application Layer



- **Internet applications**
- HTTP
- DNS
- Content delivery
- P2P

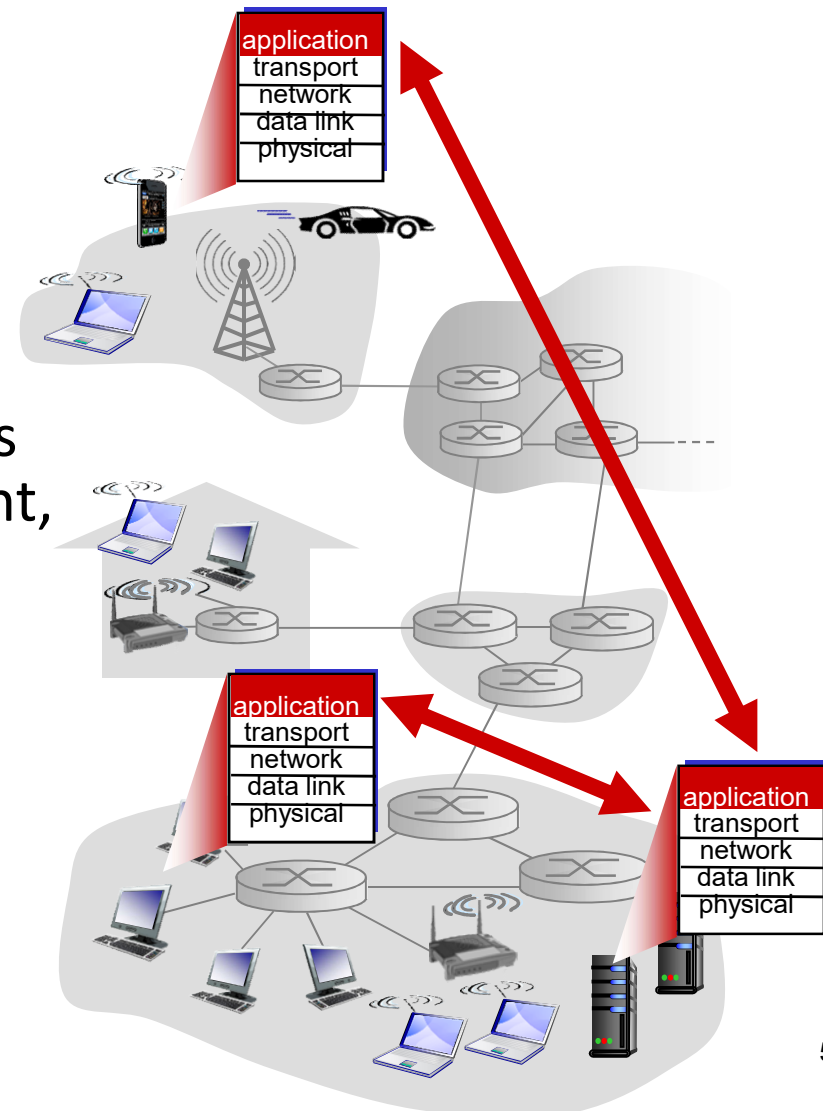
Some networked apps



Networked applications

Programs that:

- run on (different) *end systems*
- communicate over network
e.g., web server software communicates with mobile browser software
- applications only on end systems
allows for rapid app development, propagation
- two main interaction types
client – server
peer to peer




Client-server architecture

server:

- always-on host
- permanent IP address
- serves multiple clients
- data centres for scaling
- cloud computing

clients:

- communicate with server
- may be intermittently connected
- may have dynamic IP addresses
- do not communicate directly with each other




Anatomy of a datacentre



<http://www.youtube.com/watch?v=avP5d16wEp0>

P2P architecture

- *no* always-on server
 - arbitrary end systems directly communicate
 - peers request service from other peers, provide service in return to other peers
- every peer acts as client and server
- self scalability* – new peers bring new service capacity, as well as new service demands
- peers are intermittently connected and change IP addresses
- complex management
hard to provide guarantees



App-layer protocol defines



- **types of messages exchanged**
e.g., request, response
- **message syntax**
what fields in messages & how fields are delineated
- **message semantics**
meaning of information in fields
- **rules**
when and how processes send & respond to messages

The Application Layer



- Internet applications
- **HTTP**
- DNS
- Content delivery
- P2P

Web and HTTP

web page consists of *objects*

object can be HTML file, XML-Json data, js
client-side code, JPEG image, audio file,...

web page consists of *base HTML-file* which
includes *several referenced objects*

each object is addressable by a *URL*, e.g.,

`www.someschool.edu/someDept/pic.gif`

host name

path name

HTTP overview

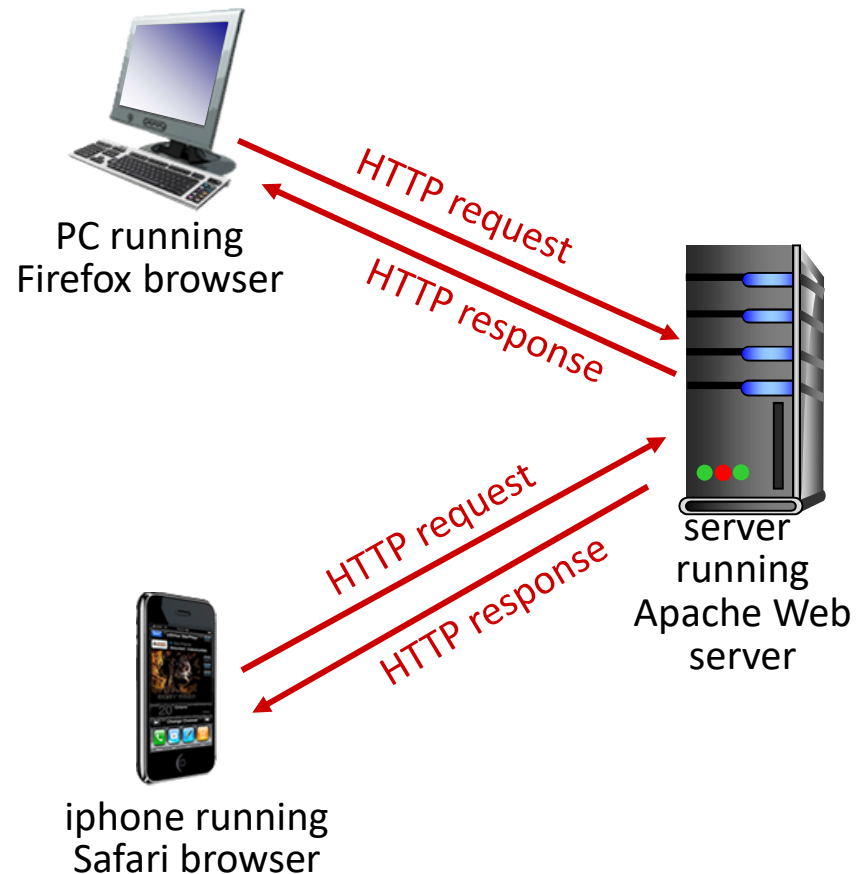
HTTP: hypertext transfer protocol

Web's application layer protocol

client/server model

client: browser that requests, receives, (using HTTP protocol) and "displays" Web objects

server: Web server sends (using HTTP protocol) objects in response to requests



HTTP overview (continued)



uses TCP:

client initiates TCP connection
(creates socket) to server,
port 80

server accepts TCP connection
from client

HTTP messages (application-
layer protocol messages)
exchanged between
browser (HTTP client) and
Web server (HTTP server)

TCP connection closed

HTTP is “stateless”

server maintains no
information about
past client requests

aside

protocols that maintain
“state” are complex!

- ❖ past history (state) must
be maintained
- ❖ if server/client crashes,
their views of “state”
may be inconsistent,
must be reconciled

Sample HTTP interaction

suppose user enters URL:

`www.someSchool.edu/someDepartment/home.index`

(contains text,
references to 10
jpeg images)

1a. HTTP client initiates TCP connection to HTTP server (process) at `www.someSchool.edu` on port 80

1b. HTTP server at host `www.someSchool.edu` waiting for TCP connection at port 80. "accepts" connection, notifying client

2. HTTP client sends HTTP *request message* (containing URL) into TCP connection socket. Message indicates that client wants object `someDepartment/home.index`

3. HTTP server receives request message, forms *response message* containing requested object, and sends message into its socket

time

Non-persistent HTTP (cont.)

4. HTTP client receives response message containing html file, displays html. Parsing html file, finds 10 referenced jpeg objects

time

5. Steps 1-4 repeated for each of 10 jpeg objects

6. HTTP server closes TCP connection.

HTTP request message

two types of HTTP messages: *request, response*

HTTP request message:

ASCII (human-readable format)

request line
(GET, POST,
HEAD commands)

header
lines

carriage return,
line feed at start
of line indicates
end of header lines

```
GET /index.html HTTP/1.1\r\n
Host: www-net.cs.umass.edu\r\n
User-Agent: Firefox/3.6.10\r\n
Accept: text/html,application/xhtml+xml\r\n
Accept-Language: en-us,en;q=0.5\r\n
Accept-Encoding: gzip,deflate\r\n
Accept-Charset: ISO-8859-1,utf-8;q=0.7\r\n
Keep-Alive: 115\r\n
Connection: keep-alive\r\n
\r\n
```

carriage return character
line-feed character

HTTP Request Methods



GET Method

Requests an object to the host

Request body is empty

Information can be encoded in the request path

Server returns the requested object

POST Method

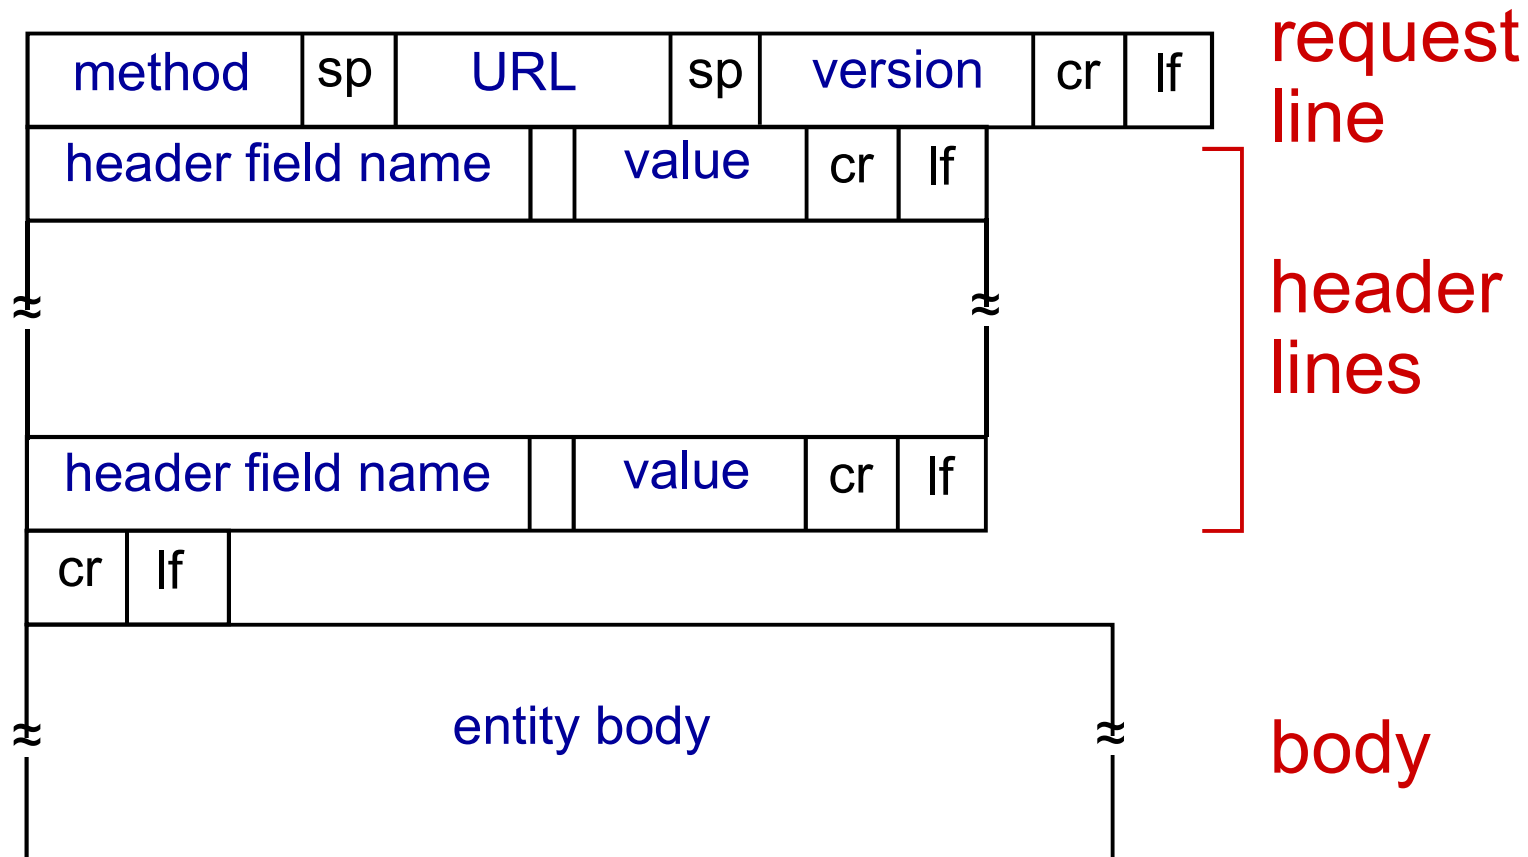
Request body contains data to be uploaded to the server

Used for uploading complex forms and data

e.g. Facebook Photo Upload

HEAD, PUT, DELETE,...

request message format



HTTP response message

status line
(protocol
status code
status phrase)

header
lines

data, e.g.,
requested
HTML file

```
HTTP/1.1 200 OK\r\n
Date: Sun, 26 Sep 2010 20:09:20 GMT\r\n
Server: Apache/2.0.52 (CentOS)\r\n
Last-Modified: Tue, 30 Oct 2007 17:00:02
GMT\r\n
ETag: "17dc6-a5c-bf716880"\r\n
Accept-Ranges: bytes\r\n
Content-Length: 2652\r\n
Keep-Alive: timeout=10, max=100\r\n
Connection: Keep-Alive\r\n
Content-Type: text/html; charset=ISO-8859-
1\r\n
\r\n
data data data data data ...
```

HTTP response status codes



- status code in 1st line in server-to-client response
- Informs about the result of the client request
- codes are predefined, numeric in the form XXX

200 OK

request succeeded, requested object embedded in message body

301 Moved Permanently

requested object moved, new location specified later in this msg
(Location:)

client can send a new request to the updated location

400 Bad Request

request msg not understood by server

404 Not Found

requested document not found on this server

505 HTTP Version Not Supported

State tracking: cookies



Problem:

http is stateless: servers do not keep track of their clients.

This complicates several tasks:

- authorization
- shopping carts
- personalized content (recommendations)
- user session state (Web e-mail)

cookies and privacy:

- ❖ cookies permit sites to learn a lot about you
- ❖ you may supply name and e-mail to sites
- ❖ this has forced current UK notification normative

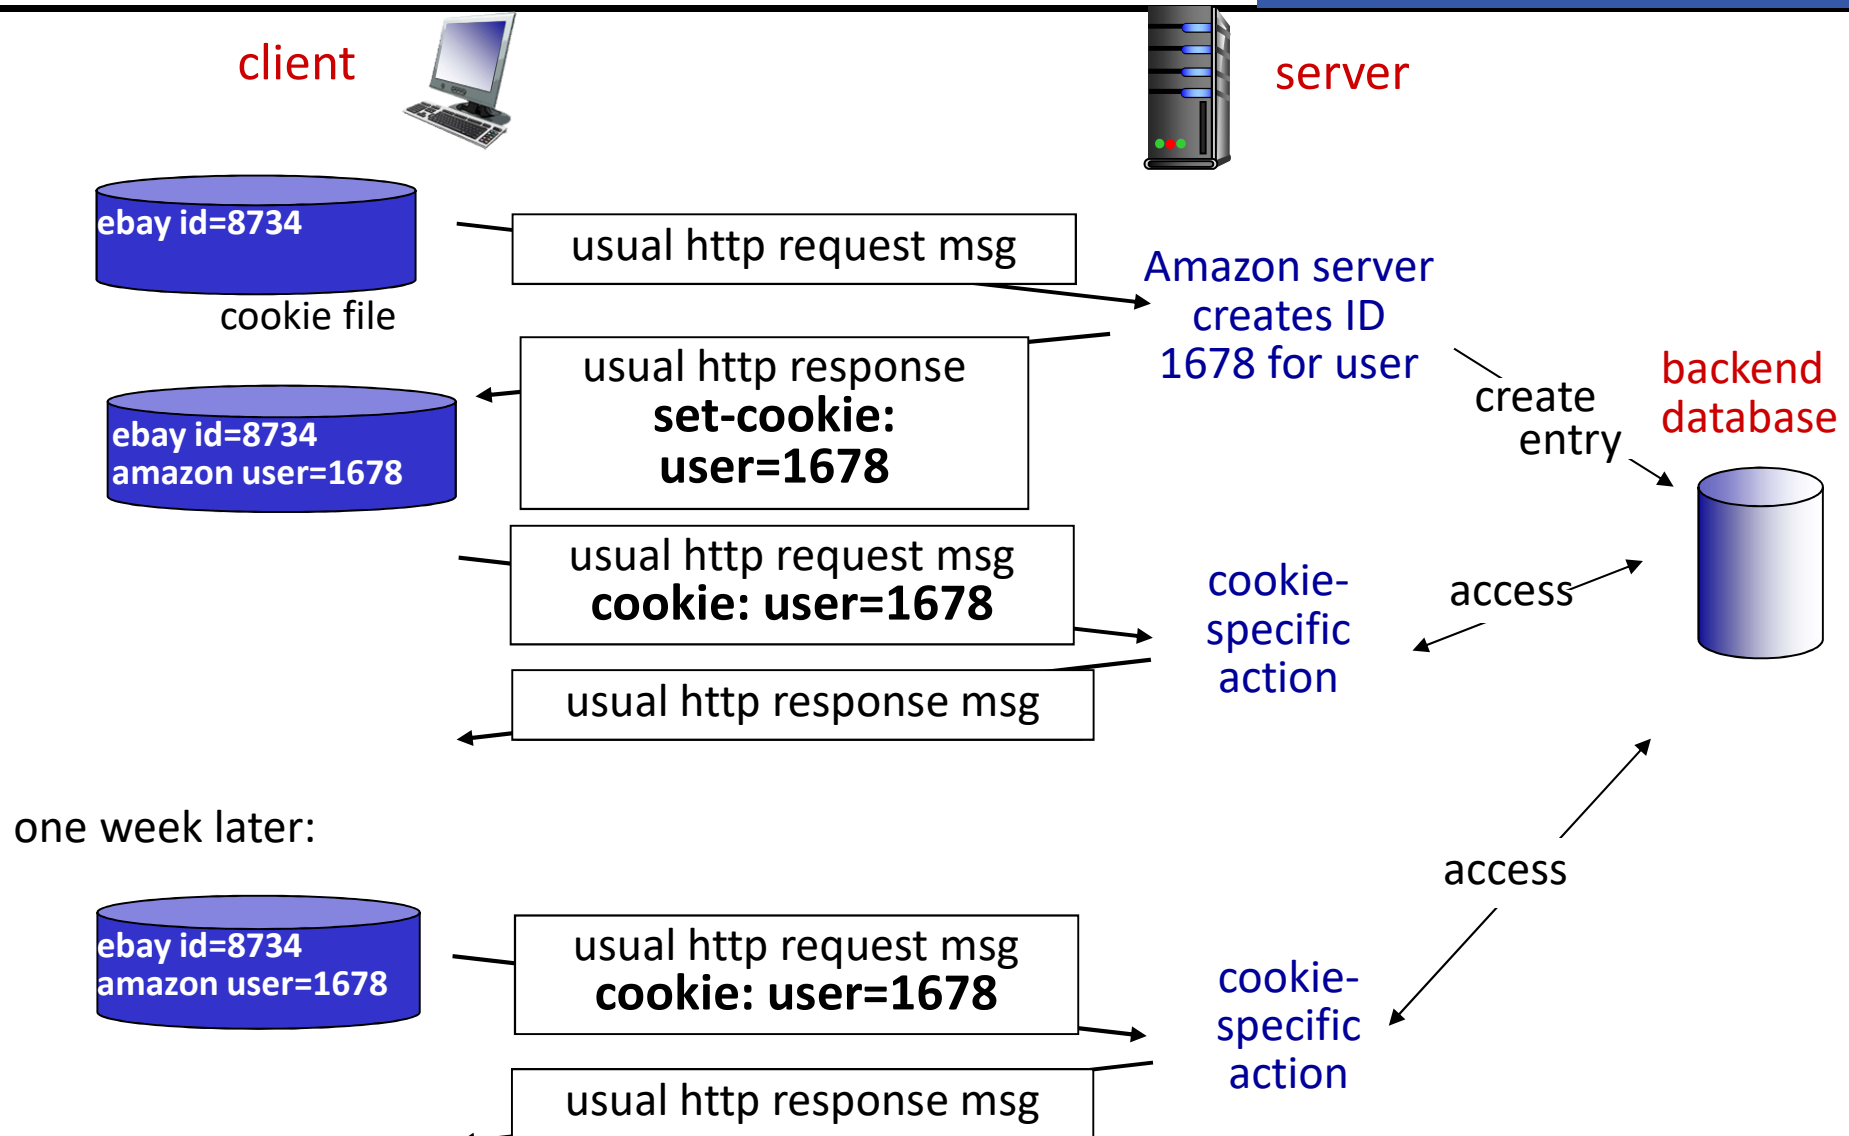
Solution:

cookies are small pieces of data, provided by the server and stored at client side.

when a browser contacts a server, a related cookie is sent to the server, after which a page can be displayed relevant to that cookie/client.

cookies belong to a domain (server), and can have a set expiration date (Expires field)

Cookies: keeping “state”




Web caches (proxy server)

cache (proxy server)
broker server to improve
efficiency

browser sends all HTTP
requests to cache

object in cache: cache
returns object

else cache requests
object from origin
server, then returns
object to client



More about Web caching



cache acts as both client
and server

server for original requesting
client

client to origin server

typically cache is
installed by ISP
(university, company,
residential ISP)

why Web caching?

reduce response time for
client request

reduce traffic on an
institution's access link

Internet dense with
caches: enables “poor”
content providers to
effectively deliver
content (so too does P2P
file sharing)

Conditional GET

Goal: don't send object if
cache has up-to-date
cached version

no object transmission delay

lower link utilization

cache: specify date of
cached copy in HTTP
request

If-modified-since: <date>

server: response contains no
object if cached copy is
up-to-date:

HTTP/1.1 304 Not Modified

client



HTTP request msg
If-modified-since: <date>

server



HTTP response
**HTTP/1.1
304 Not Modified**

object
not
modified
before
<date>



HTTP request msg
If-modified-since: <date>

HTTP response
**HTTP/1.1 200 OK
<data>**

object
modified
after
<date>

RESTful applications



REpresentational State Transfer. Roy Fielding,
2000

Client-server architecture for distributed
applications

Information is abstracted as resources

- Resources are stored at the server, they are stateful

- Resources have identity (can be referenced)

- Uniform resource access interface

- Client and server exchange resource representations
(eg XML or JSON format)

Http is perfectly suited for RESTful services

REST Resource Identification



URI: Uniform Resource Identifier

`scheme://host:port/path?queryString#fragment`

A resource is identified by its URI

URLs act as templates

Part of the address represents the type of the resource, and the other part the individual resource

`http://phones.com/whitepages/{userPhone}`

Accept: header specifies the accepted data representation formats (json nowadays)

REST HTTP methods

Clients interact with resources by invoking HTTP methods

Request method dictates what the client wants to do

Method	Meaning	Request Body	Response Body
GET	<i>Retrieve</i> resource representation	Empty	Resource representation
DELETE	<i>Delete</i> the resource	Empty	Empty (or status)
PUT	<i>Create new</i> resource	Proposed representation	Resource representation
POST	Edit/ <i>Update</i> resource	Proposed representation	Resource representation

Resource representation (json)



Json has become the de-facto standard

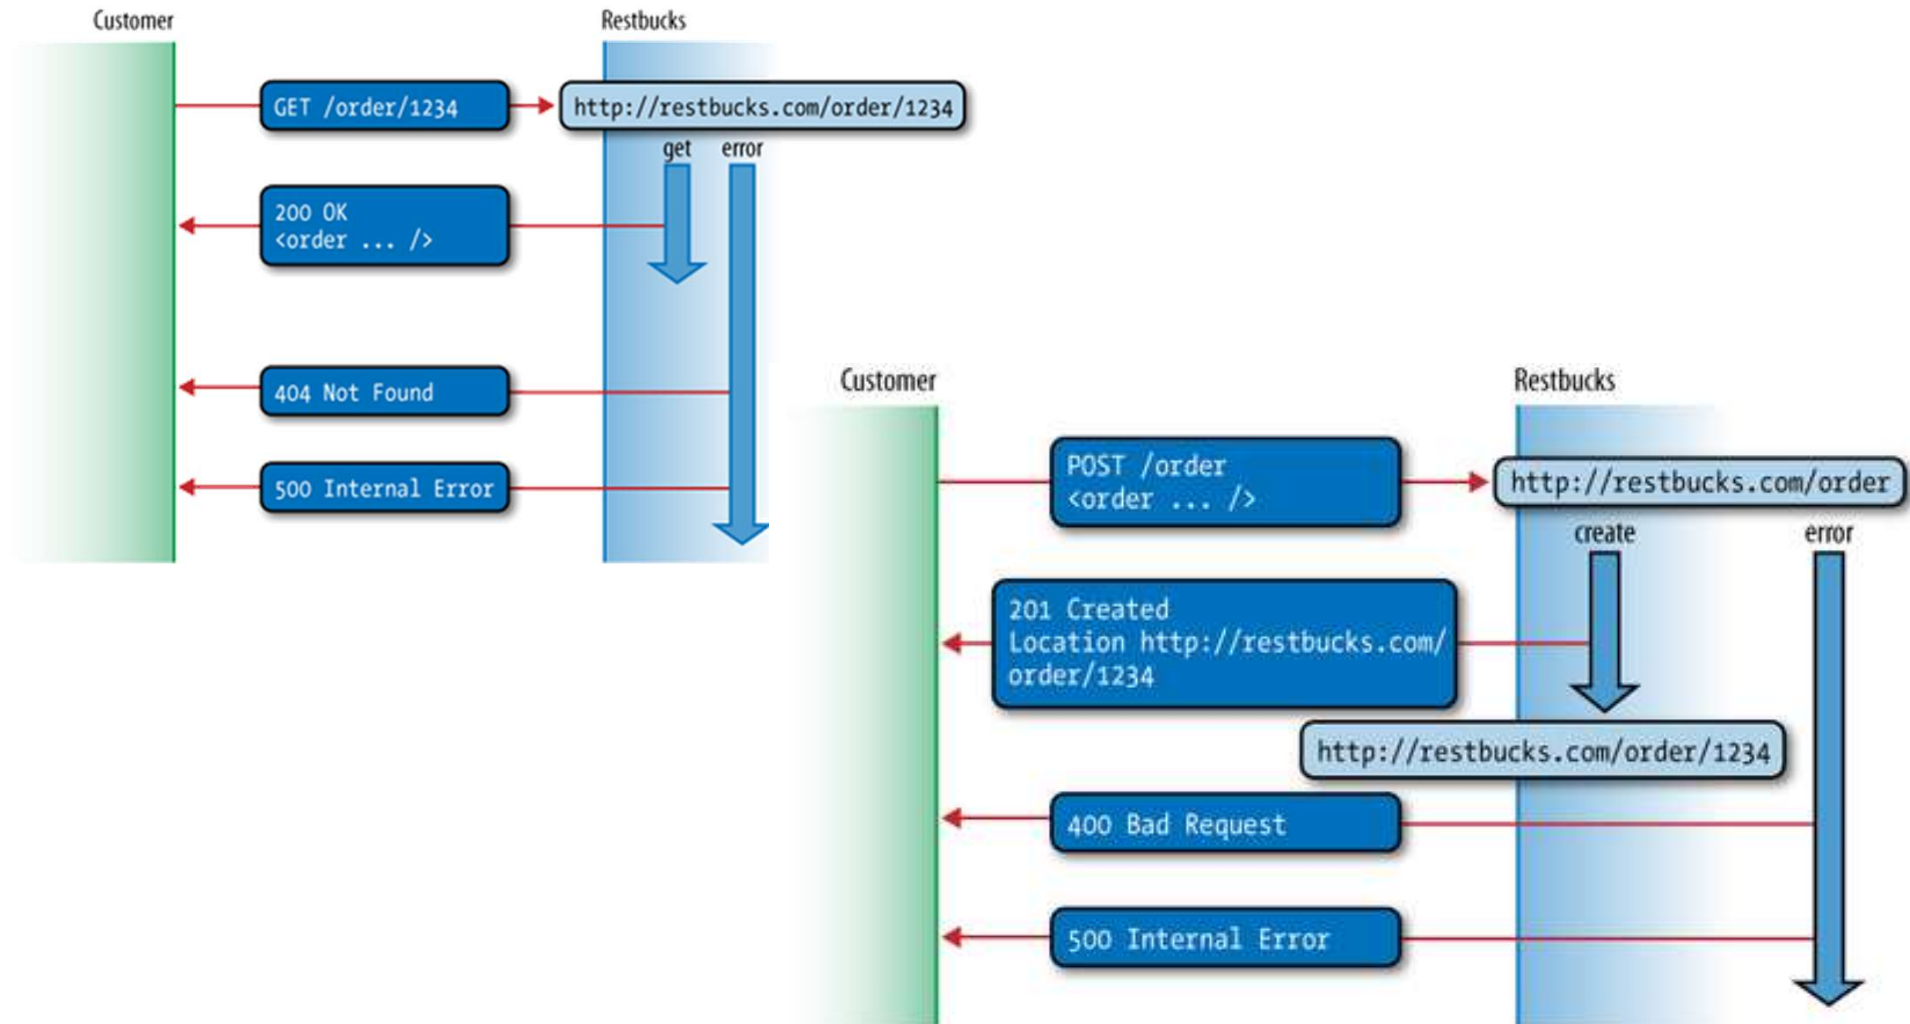
Originally part of javascript, text-like format for describing tree-like structures

```
{ "id":"26007494656", "user_id":"23161357",  
  "game_id":"417752",  
  "community_ids":[ "5181e78f-2280-42a6-873d58e25a7c313",  
    "848d95be-90b3-44a5-b143-6e373754c382", "fd0eab99-832a-4d7e-  
    8cc0-04d73deb2e54" ],  
  "type":"live",  
  "title":"Hey Guys, It's Monday - Twitter: @Lirik",  
  "viewer_count":32575,  
  "started_at":"2017-08-14T16:08:32Z",  
  "language":"en"}, }
```


Example REST API for orders

Method	Meaning
GET <i>/order/{orderId}</i>	Requests the current status of the order whose identifier is <i>orderId</i>
DELETE <i>/order/{orderId}</i>	Removes the order with the provided <i>orderId</i>
PUT <i>/order/{orderId}</i>	Updates the order with the provided <i>orderId</i> . Request body includes updated representation.
POST <i>/order</i>	Creates new order. Response message contains the id for the newly created resource

REST service interactions



Exchange of HTTP messages

```
POST /order HTTP/1.1
Host: restbucks.com
Content-Type: application/xml
Content-Length: 239
{  "order": {
    "location": "takeAway",
    "items": {
      "item": {
        "name": "latte",
        "quantity": "1",
        "milk": "whole",
        "size": "small"
      }
    }
  }
}
```

```
HTTP/1.1 201 Created
Content-Length: 267
Content-Type: application/xml
Date: Wed, 26 Sep 2017 08:45:03
Location:
http://restbucks.com/order/1234
{ "order" : { ...
```

```
GET /order/1234 HTTP/1.1
Host: restbucks.com
```

```
HTTP/1.1 200 OK
Content-Length: 241
Content-Type: application/xml
Date: Wed, 26 Sep 2017 09:02:10
{  "order": {
    "location": "takeAway",
    "items": {
      "item": {
        "name": "latte",
        "quantity": "1",
        "milk": "whole",
        "size": "small"
      }
    }
    "status": delivered }
}
```

Using REST in client applications



It is possible to create a REST client by creating manually http messages, and sending them

... but popularity of REST has led to a myriad of high-level libraries for every framework

In Java, Jackson eases json management

```
public Book[] getAllBooks() throws  
Exception {  
    HttpResponse<Book[]> response =  
    Unirest.get(URI_BOOK).asObject(Book[].cl  
ass);  
    Book[] books = response.getBody();  
    return books;  
}
```

<https://howtoprogram.xyz/2016/07/27/java-rest-client-using-unirest-java-api/>

The Application Layer



- Internet applications
- HTTP
- **DNS**
- Content delivery
- P2P

DNS: domain name system



- *People*: many identifiers:
name, passport #
- *Internet hosts, routers*:
 - IP address (32 bit) -
used for addressing
datagrams
 - “name”, e.g.,
www.yahoo.com -
used by humans

Q: how to map between IP
address and name, and
vice versa?

Domain Name System:

- *distributed database*
implemented in hierarchy of
many *name servers*
- *application-layer protocol*:
hosts, name servers
communicate to *resolve*
names (address/name
translation)
 - note: core Internet function,
implemented as application-
layer protocol
 - complexity at network's
“edge”

DNS: services, structure



DNS services


- hostname to IP address translation
- host aliasing
 - canonical, alias names
- mail server aliasing
- load distribution
 - replicated Web servers:
 - many IP addresses
 - correspond to one name

Why not centralize DNS?

- single point of failure
- traffic volume
- distant centralized database
- maintenance

A: doesn't scale!

DNS: a distributed, hierarchical database

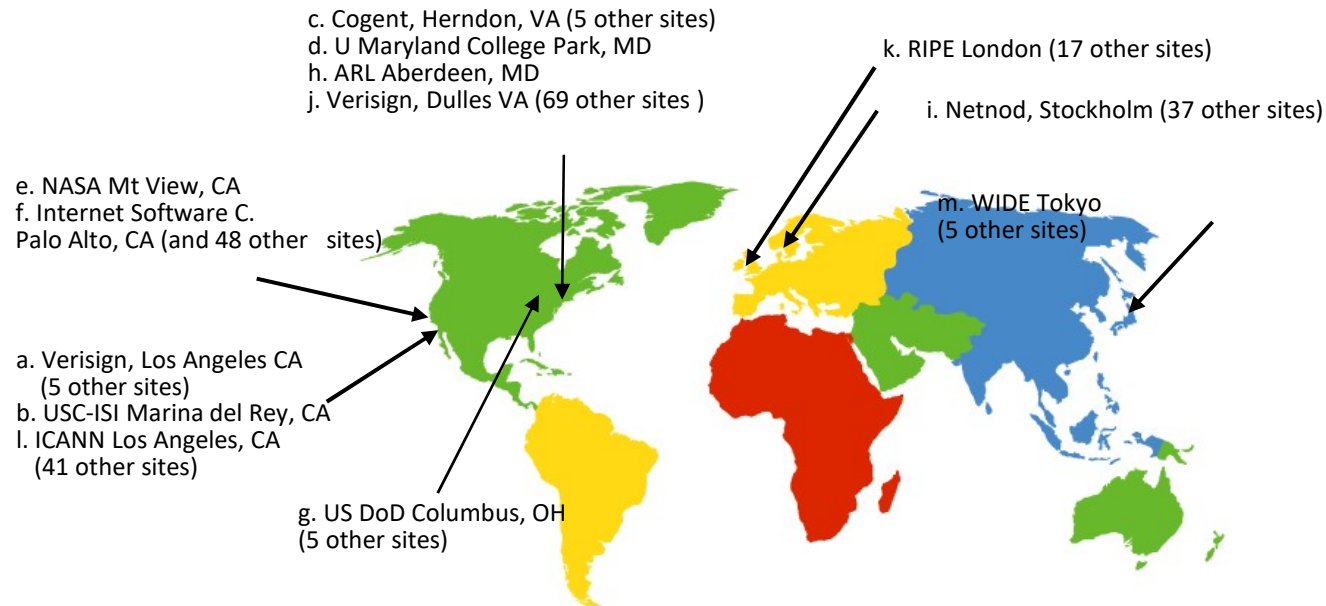


Client wants IP for www.amazon.com; 1st approx:

- client queries root server to find com DNS server
- client queries .com DNS server to get amazon.com DNS server
- client queries amazon.com DNS server to get IP address for www.amazon.com

DNS: root name servers

- Contacted by local resolver that cannot resolve name
- Root name server:
 - contacts authoritative name server if name mapping not known
 - gets mapping
 - returns mapping to local name server



- *Top-level domain (TLD) servers:*
 - responsible for com, org, net, edu, aero, jobs, museums, and all top-level country domains, e.g.: uk, fr, ca, jp
 - Network Solutions maintains servers for .com TLD
 - Educause for .edu TLD
- *Authoritative DNS servers:*
 - Organization's own DNS server(s), providing authoritative hostname to IP mappings for organization's named hosts
 - can be maintained by organization or service provider

Local DNS resolver

- Client, not part of DNS hierarchy
- Each ISP (residential ISP, company, university) has one
 - also called “default resolver”
- When host makes DNS query, query is sent to its local DNS resolver
 - has local cache of recent name-to-address translation pairs (but may be out of date!)
 - acts as a proxy, forwards query into DNS hierarchy
- 3rd party resolvers
 - OpenDNS
 - GoogleDNS

DNS name resolution example

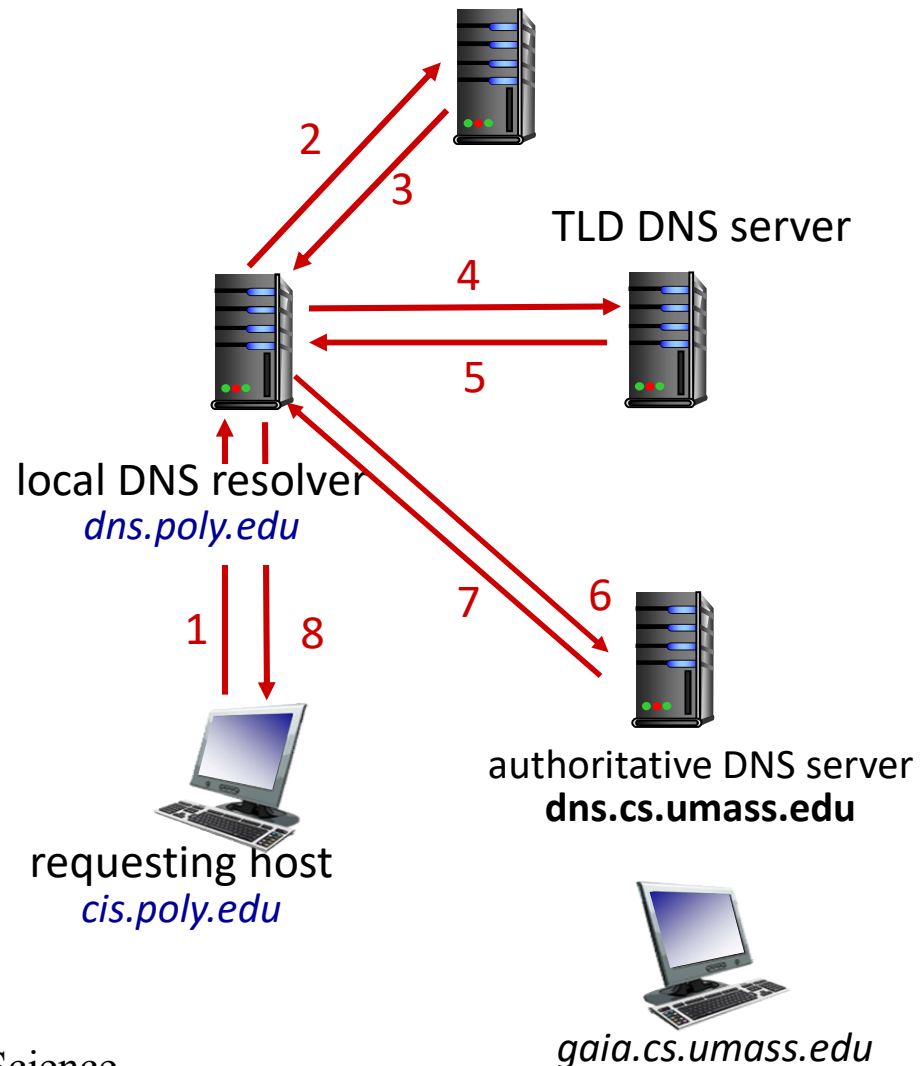


root DNS server

host at cis.poly.edu wants
IP address for
gaia.cs.umass.edu

Iterative query:

- Contacted server replies with name of server to contact (referral) = “I don’t know this name, but ask this server”
- Resolver iterates until it finds the answer



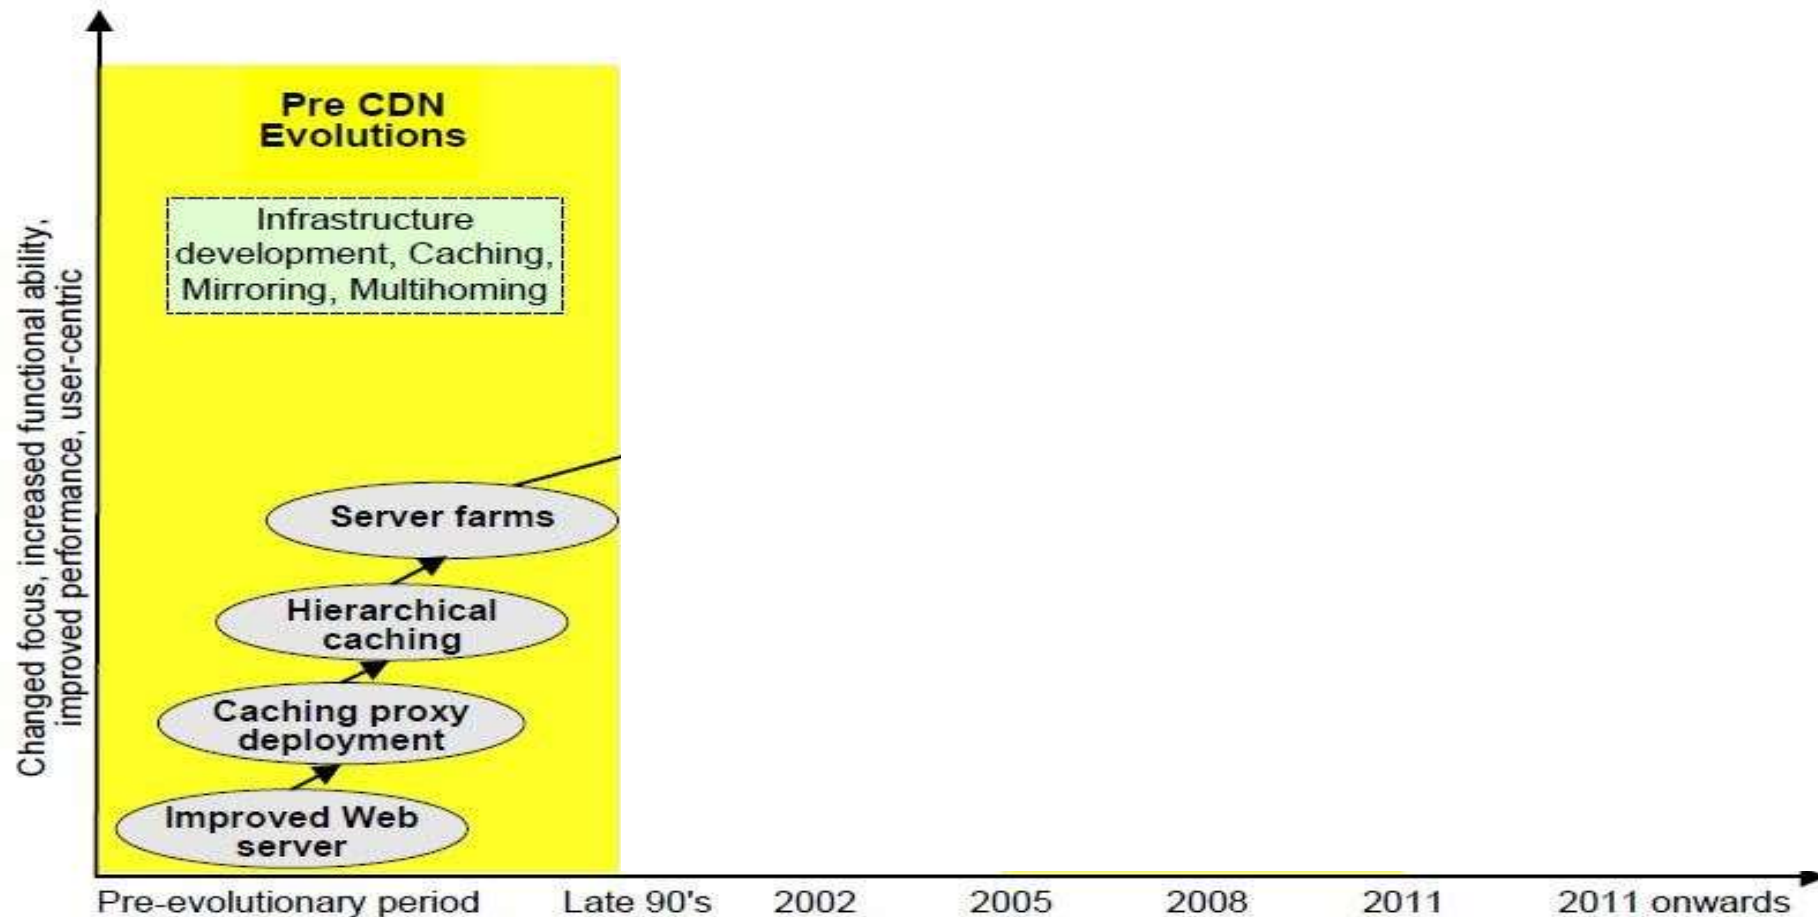
DNS caching

- Once (any) name server/resolver learns mapping, it *caches* mapping
 - cache entries timeout (disappear) after some time (TTL)
 - TLD servers typically cached in local name servers
 - => root name servers not often visited
- Cached entries may be *out-of-date* (best effort name-to-address translation!)
 - if name host changes IP address, may not be known Internet-wide until all TTLs expire

Agenda

- Internet applications
- HTTP
- DNS
- **Content delivery**
- P2P

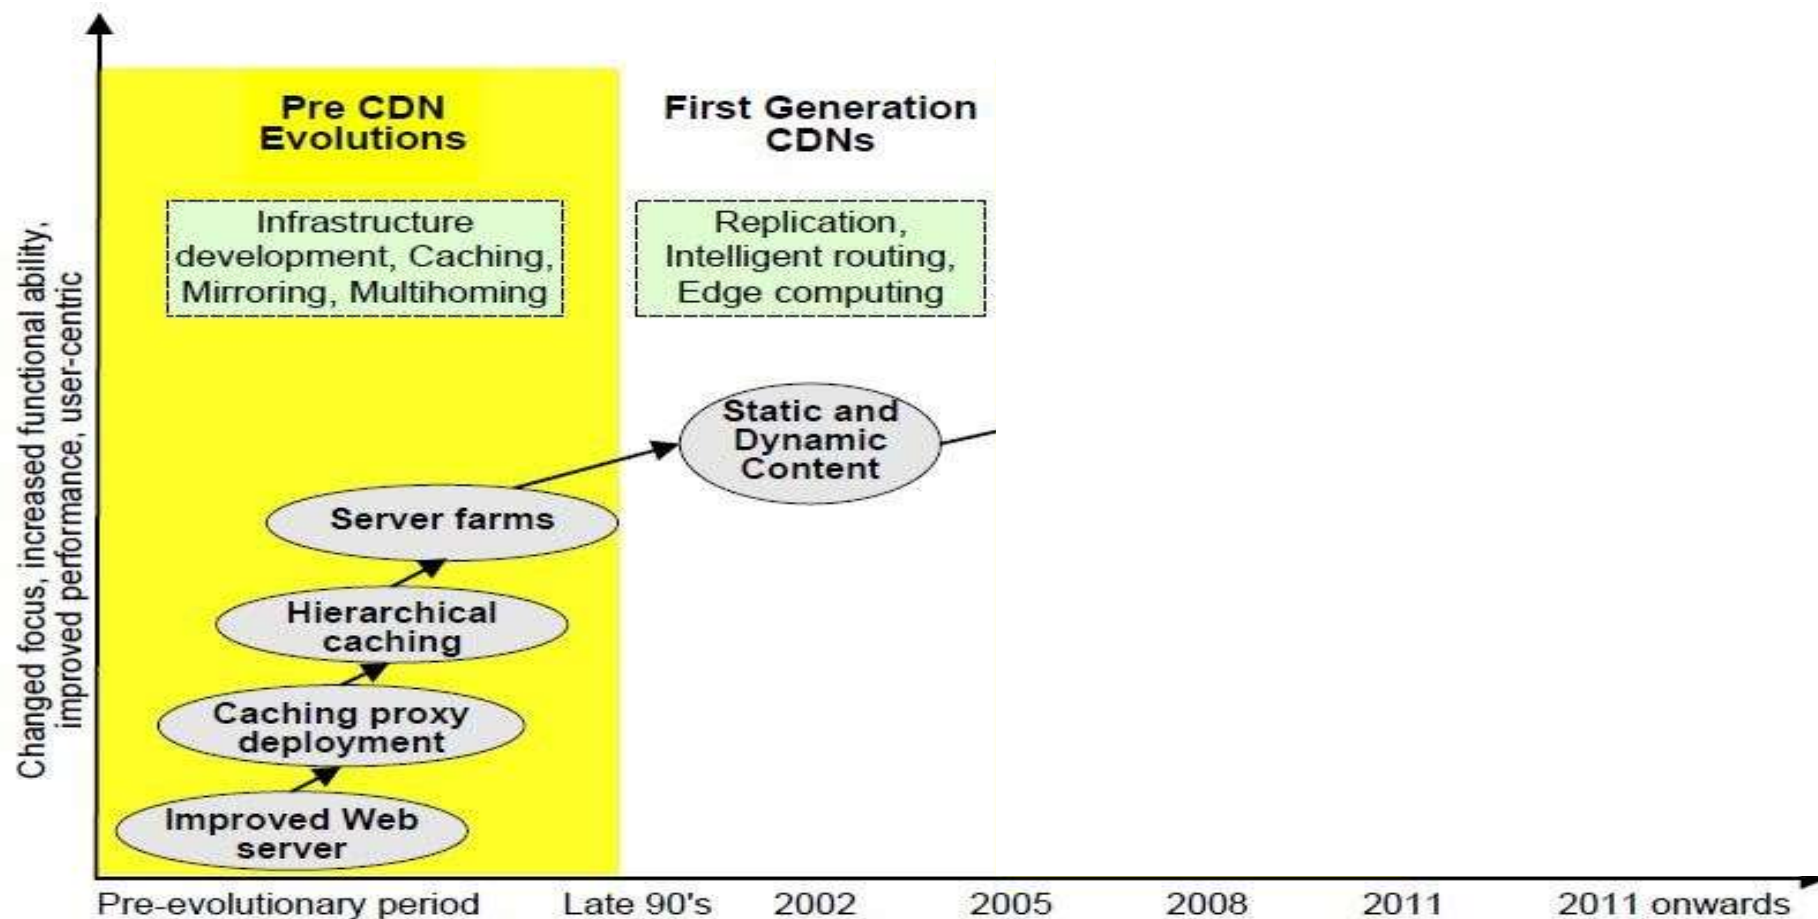
The Early Web



Pathan Mukaddim. Ongoing Trends and Future Directions in Content Delivery Networks (CDNs). Available online from: <http://amkpathan.wordpress.com/article/ongoing-trends-and-future-directions-in-3uxfz2buz8z1w-2/>

School of Electronic Engineering and Computer Science

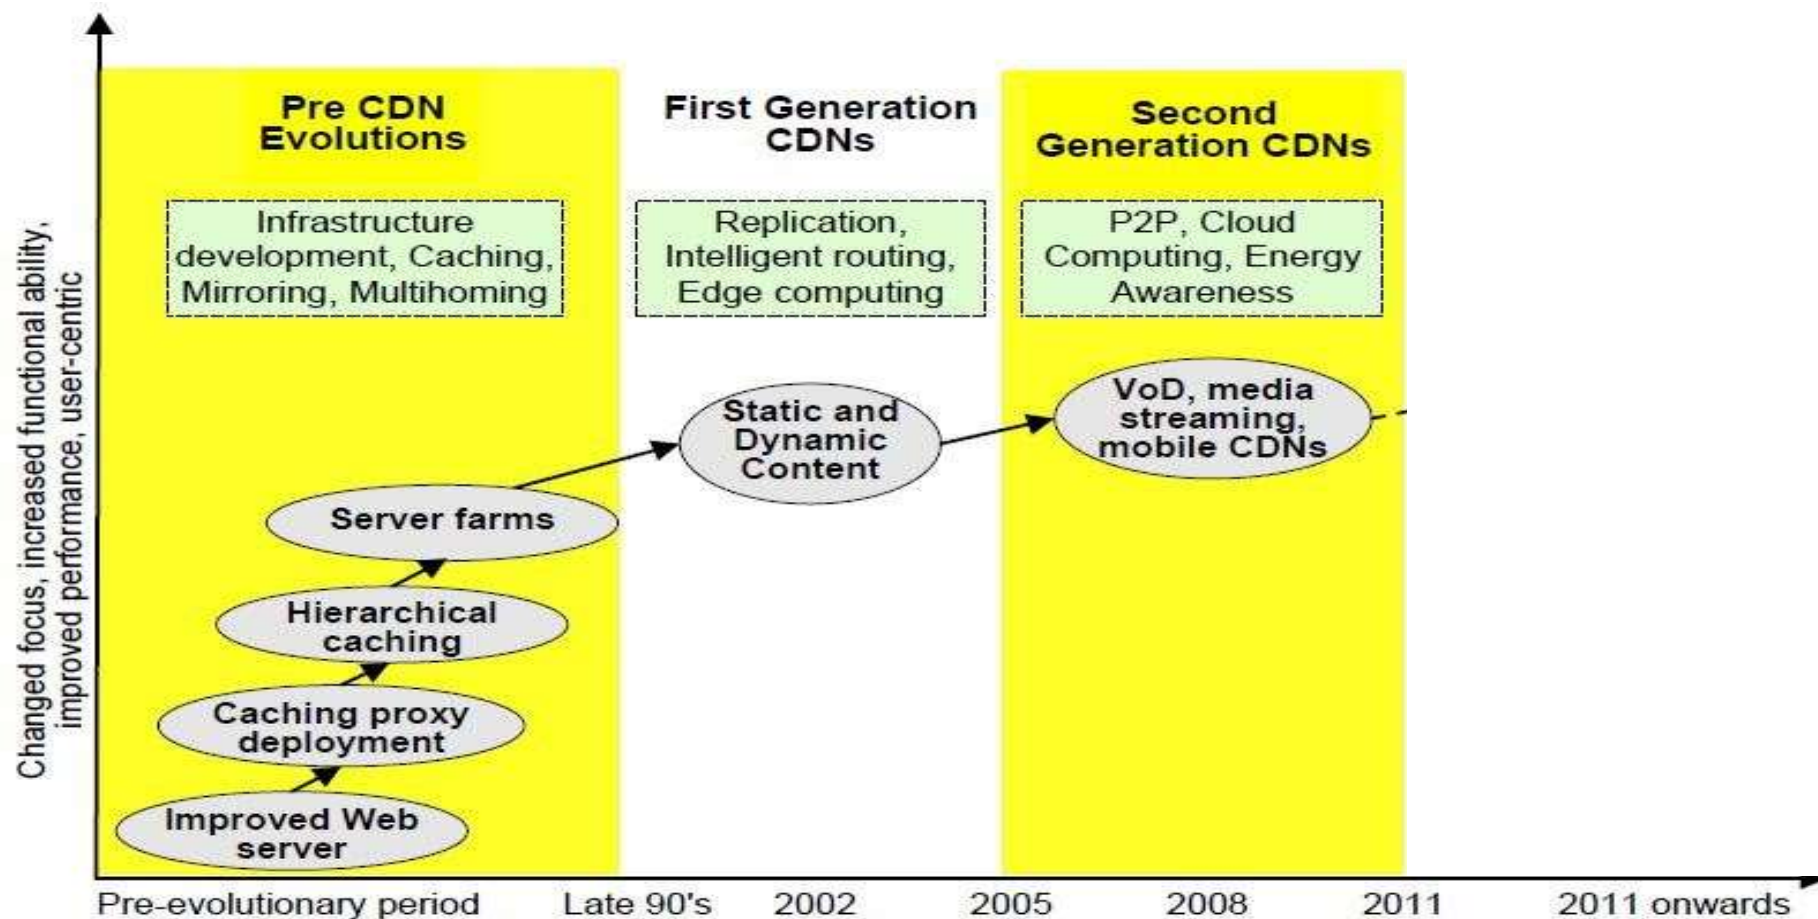
CDNs 1.0



Pathan Mukaddim. Ongoing Trends and Future Directions in Content Delivery Networks (CDNs). Available online from: <http://amkpathan.wordpress.com/article/ongoing-trends-and-future-directions-in-3uxfz2buz8z1w-2/>

School of Electronic Engineering and Computer Science

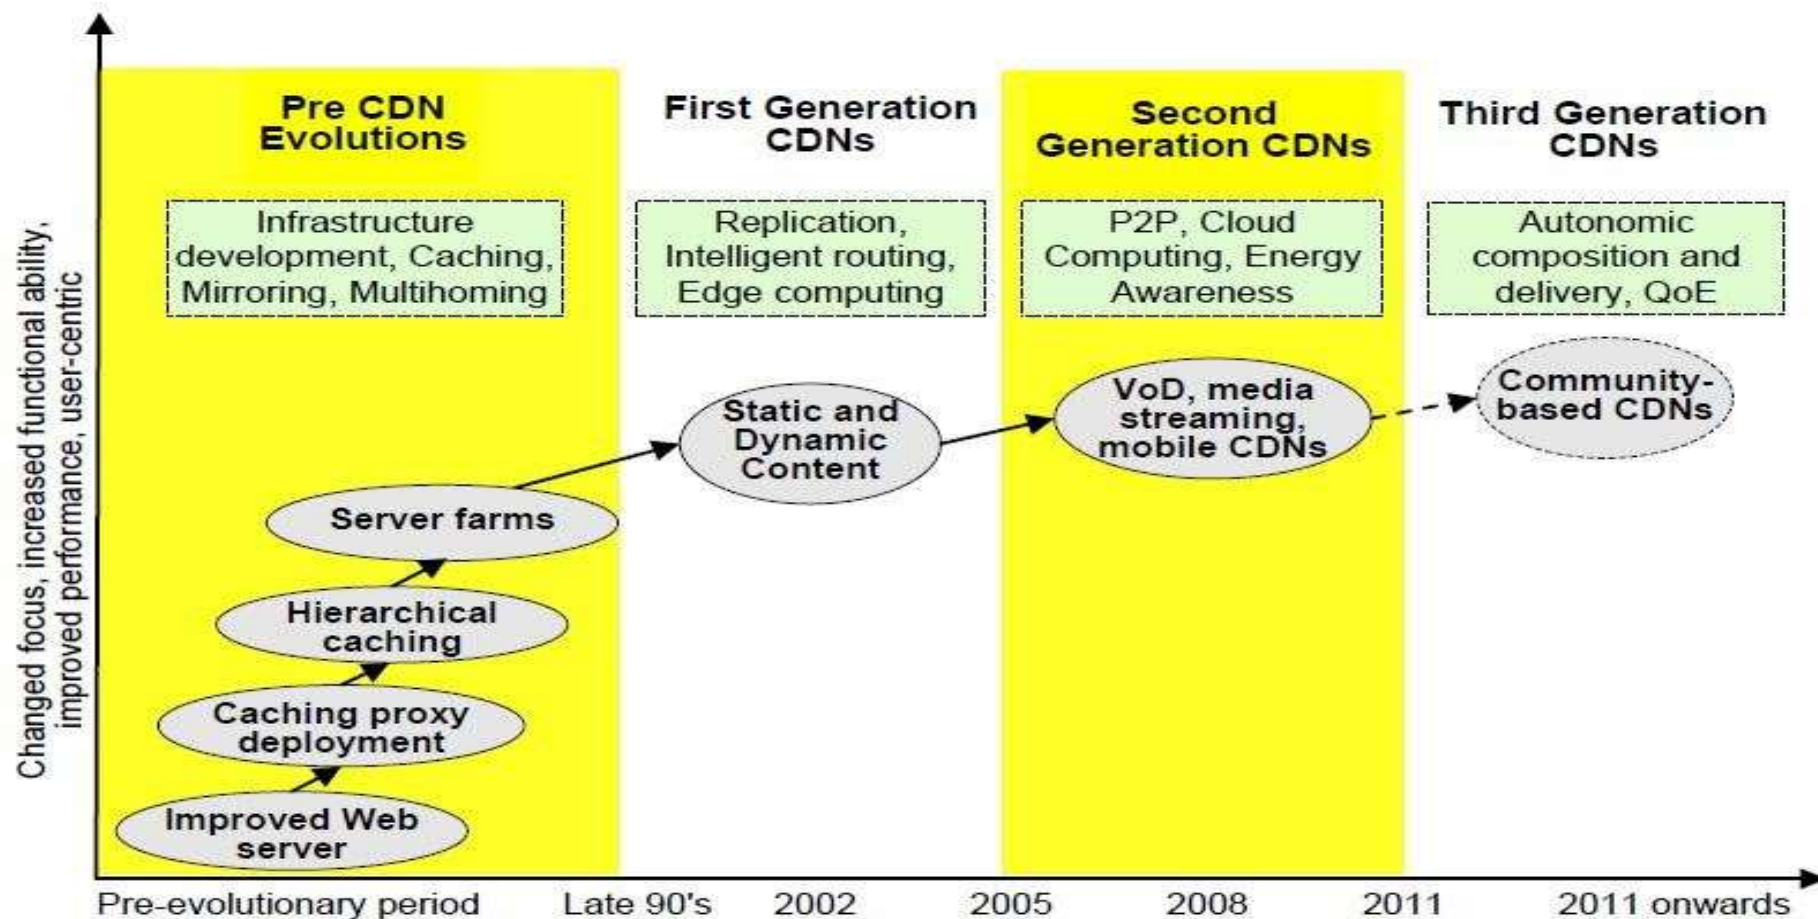
CDNs 2.0



Pathan Mukaddim. Ongoing Trends and Future Directions in Content Delivery Networks (CDNs). Available online from: <http://amkpathan.wordpress.com/article/ongoing-trends-and-future-directions-in-3uxfz2buz8z1w-2/>

School of Electronic Engineering and Computer Science

Autonomic CDNs

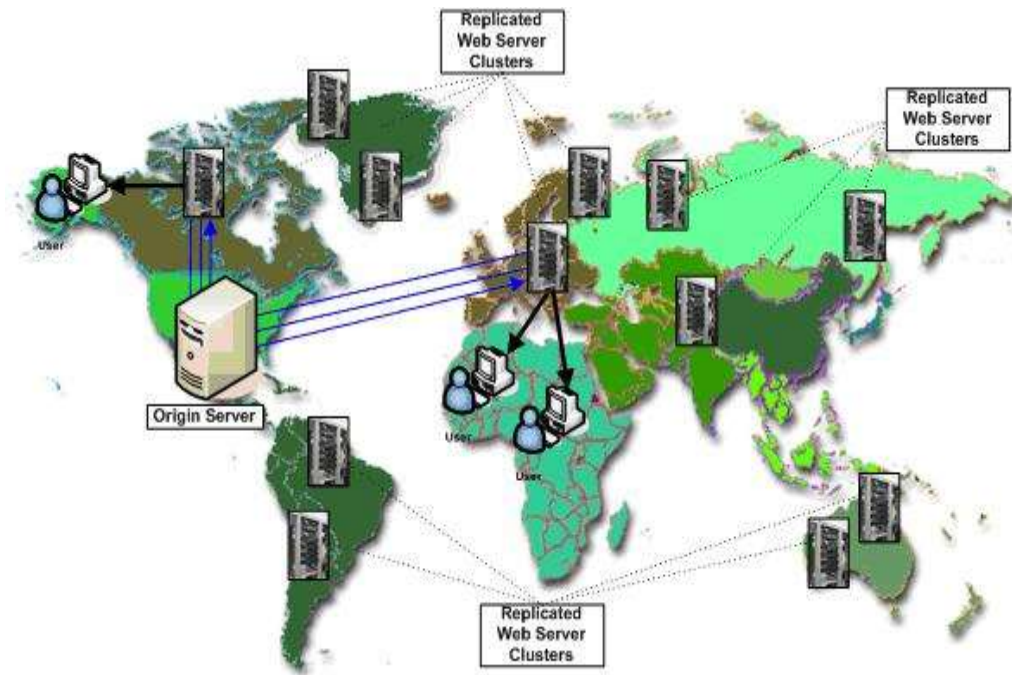


Pathan Mukaddim. Ongoing Trends and Future Directions in Content Delivery Networks (CDNs). Available online from: <http://amkpathan.wordpress.com/article/ongoing-trends-and-future-directions-in-3uxfz2buz8z1w-2/>


School of Electronic Engineering and Computer Science

The CDN Monsters

- Servers deployed around the world
- Serve the same content from multiple (all) locations
- Examples:
 - Akamai
 - Google
 - Netflix
 - Limelight
 - Amazon CloudFront
 - Level 3
 - Windows Azure



HTTP server selection



CDNs and CNAME

- DNS redirection: CNAME = canonical name
- Popular for CDNs, e.g., Akamai

```
; <<>> DiG 9.8.3-P1 <<>> www.whitehouse.gov
```

```
;; global options: +cmd
```

```
;; Got answer:
```

```
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 11921
```

```
;; flags: qr rd ra; QUERY: 1, ANSWER: 5, AUTHORITY: 0, ADDITIONAL: 0
```

```
;; QUESTION SECTION:
```

```
;www.whitehouse.gov.          IN      A
```

```
;; ANSWER SECTION:
```

```
www.whitehouse.gov.          3596    IN      CNAME     www.whitehouse.gov.edgesuite.net.
```

```
www.whitehouse.gov.edgesuite.net. 896    IN      CNAME     www.eop-edge-lb.akadns.net.
```

```
www.eop-edge-lb.akadns.net. 300     IN      CNAME     a1128.dsch.akamai.net.
```

```
a1128.dsch.akamai.net.       20      IN      A          195.59.54.235
```

```
a1128.dsch.akamai.net.       20      IN      A          195.59.126.161
```

```
;; Query time: 43 msec
```

```
;; SERVER: 138.37.0.88#53(138.37.0.88)
```

```
;; WHEN: Fri Oct 10 10:49:15 2014
```

```
;; MSG SIZE rcvd: 183
```

World data centers



<http://www.datacentermap.com/>

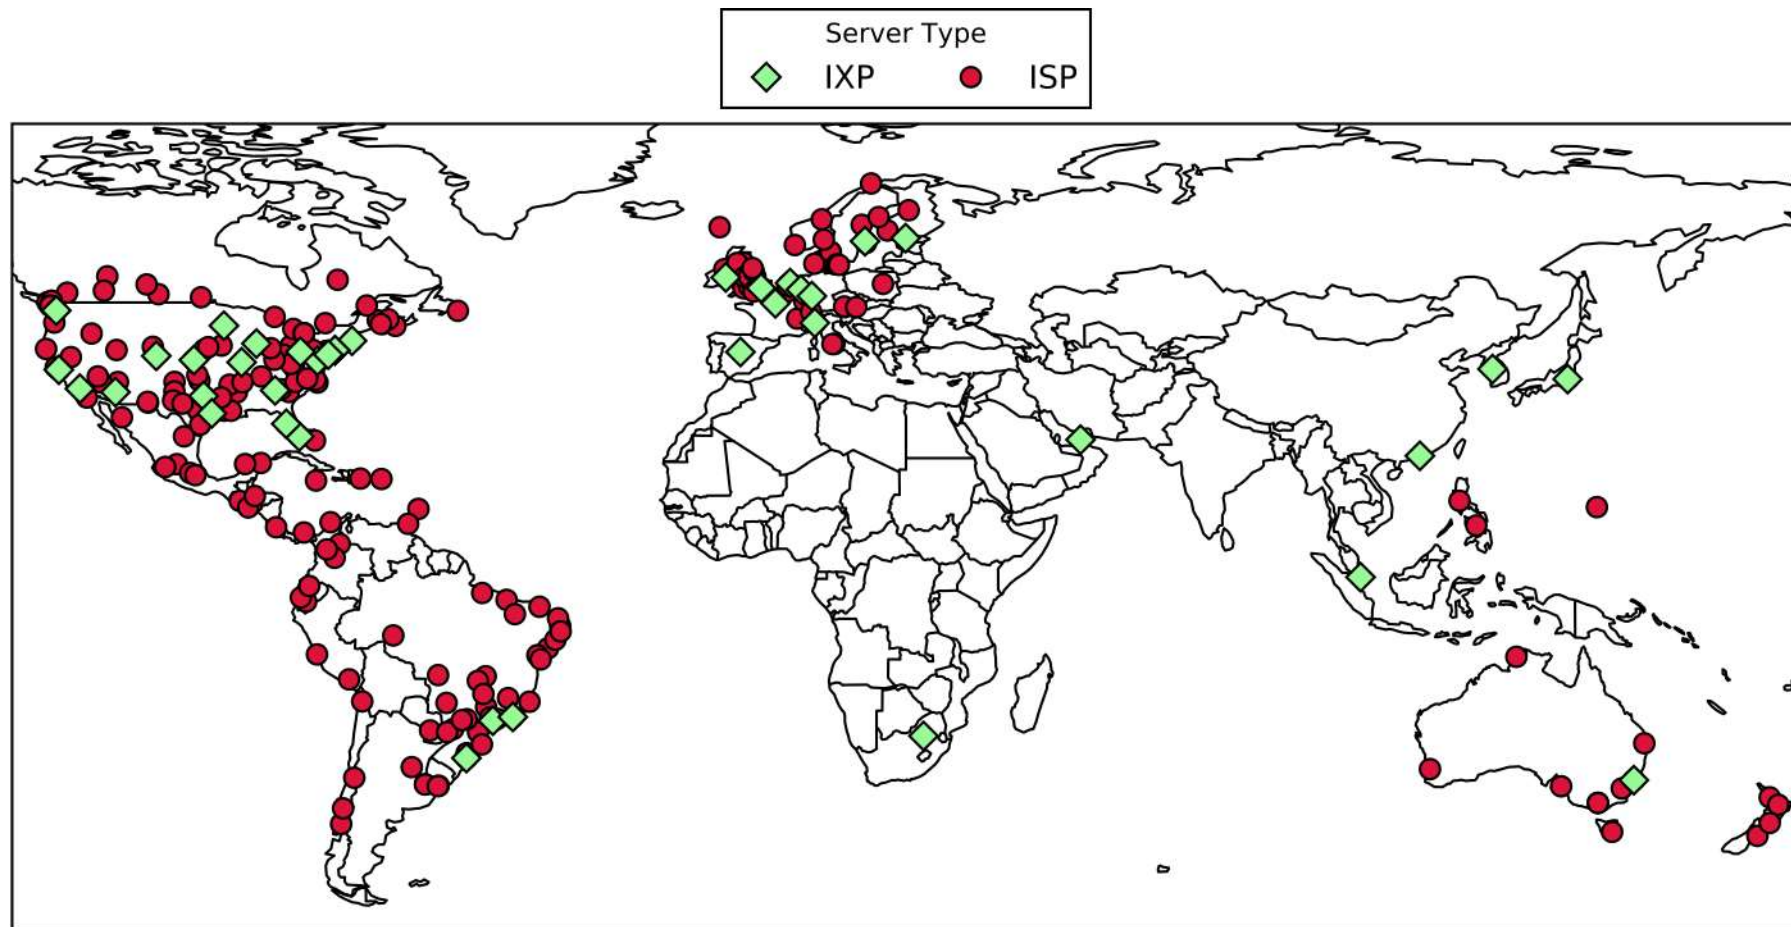
School of Electronic Engineering and Computer Science

Google data centers



<http://royal.pingdom.com/2008/04/11/map-of-all-google-data-center-locations/>
School of Electronic Engineering and Computer Science

Netflix























Where is popular Web content hosted?

- California
- China 2nd
- USA: 9 among top 20
- Other developed countries: limited own content

Rank	Country	Potential	Normalized potential
1	USA (CA)	0.254	0.108
2	China	0.128	0.107
3	USA (TX)	0.190	0.061
4	Germany	0.183	0.058
5	Japan	0.163	0.051
6	France	0.146	0.034
7	Great Britain	0.157	0.030
8	Netherlands	0.144	0.029
9	USA (WA)	0.135	0.027
10	USA (unknown)	0.164	0.027
11	Russia	0.038	0.027
12	USA (NY)	0.130	0.026
13	Italy	0.122	0.018
14	USA (NJ)	0.125	0.016
15	Canada	0.028	0.015
16	USA (IL)	0.116	0.014
17	Australia	0.118	0.013
18	Spain	0.116	0.013
19	USA (UT)	0.111	0.012
20	USA (CO)	0.113	0.012


Top western Web CDNs

Rank	# hostnames	# ASes	# prefixes	owner	content mix
1	476	79	294	Akamai	
2	161	70	216	Akamai	
3	108	1	45	Google	
4	70	35	137	Akamai	
5	70	1	45	Google	
6	57	6	15	Limelight	
7	57	1	1	ThePlanet	
8	53	1	1	ThePlanet	
9	49	34	123	Akamai	
10	34	1	2	Skyrock	
11	29	6	17	Cotendo	
12	28	4	5	Wordpress	
13	27	6	21	Footprint	
14	26	1	1	Ravand	
15	23	1	1	Xanga	
16	22	1	4	Edgecast	
17	22	1	1	ThePlanet	
18	21	1	1	ivwbox.de	
19	21	1	5	AOL	
20	20	1	1	Leaseweb	

 only on TOP,
  both on TOP and EMBEDDED,
  only on EMBEDDED,
  TAIL.


Collaborative content delivery

- Importance of stakeholders
- Spectrum in the solution space



CDN 2.0

- Hybrid infrastructures: Akamai, PPTV
- Meta-CDNs, e.g., Conviva
- Virtual CDNs through ISP micro-datacenters





The Application Layer



- Internet applications
- HTTP
- DNS
- Content delivery
- **P2P**

P2P

- Alternative to client-server model
- Each node has the same role
- Structured or unstructured
- Often built on DHT
 - Content stored as *<hash,value pair>*
 - Locating the data based on hash
 - Content replicated across P2P network
- Examples:
 - Gnutella
 - eDonkey
 - Bittorrent




P2P: Server selection

Seed:

Originator of the file (has the whole file)

Tracker: keeps track of where file copies reside on peer machines



Swarm: peers that have at least one chunk

Client: peer who wants the file

P2P: pro's and con's

- Pro's:
 - ✓ Easy to share content
 - ✓ Anonymous thanks to DHT
 - ✓ Storage scales
- Con's
 - Reliability (nodes uptime), e.g., skype
 - Lack of visibility of content (copyrighted material)
 - P2P structure dictates where content is replicated

The “P2P CDN”

#	Country	U. Peers
1	US	1379462
2	ES	887480
3	GB	800308
4	CA	551820
5	IN	514246
6	AU	322009
7	BR	318294
8	IT	295339
9	PL	288780
10	PT	220124

TOP 10 ISPs (BT VIDEO USER)			
	AS#	Peers	AS Name-Internet Service Provider
1	3352	165469	TELEFONICA-DATA-ESPANA(TDE)
2	3662	129047	DNEO-OSP7-COMCAST CABLE
3	6461	127297	MFNX MFN-METROMEIDA FIBER
4	2119	113597	TELENOR-NEXTEL T.NET
5	19262	101390	VZGNI-TRANSIT-Verizon ISP
6	3301	97658	TELIANET-SWEDEN TELIANET
7	3462	96564	HINET-DATA CBG
8	4134	87392	CHINANET-BACKBONE
9	6327	86964	SHAW-SHAW COMMUNICATION
10	174	74453	COGENT COGENT/PSI

P2P popularity

