

**ECS518U - Operating Systems**  
**Week 6**

# **Memory Management: Address Spaces & Paging**

Tassos Tombros

# Outline

- **The problem** that memory management tries to solve
  - Code relocation
  - Efficient memory use
- **Solutions**
  - Segments
  - Pages
- **Virtual memory**
- **Locality and page faults**
- **OS design** issues
  - Page replacement, ...
- **Reading:**
  - **Stalling:** Chapter 7 & Chapter 8 (8.1, 8.2)
  - **Tanenbaum:** Chapter 3

Today

Week 8

# Things you will learn today

- Why is Memory Management necessary, **what problems does it solve?**
- Different address spaces
  - The logical address space
  - The physical address space
- **Solutions** to the Memory Management problem:
  - Fixed & Dynamic **Partitioning**
  - **Segmentation**
  - **Paging**

# IMPORTANT: LOOKING AHEAD

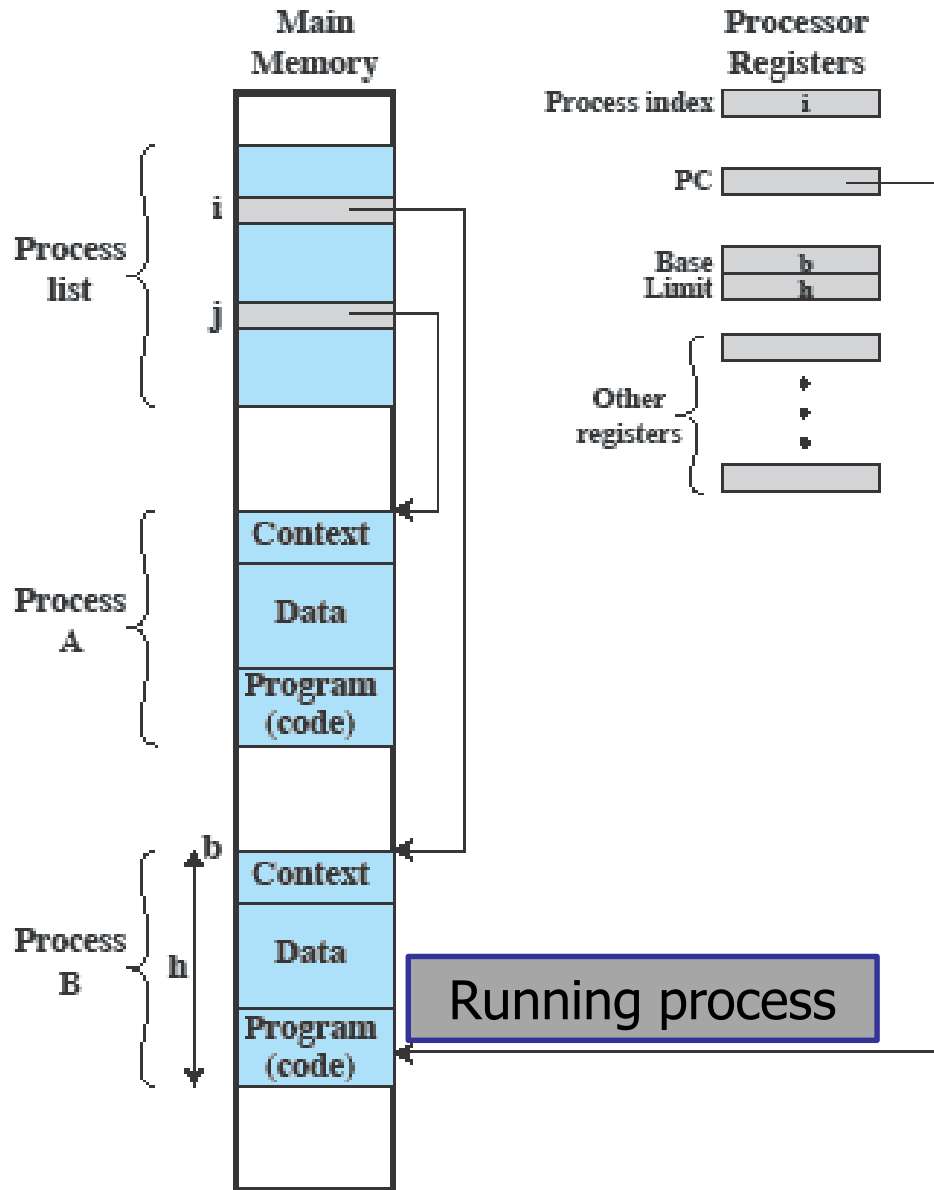
- **WEEK 7**

- **1<sup>ST</sup> MCQ Test will be released**
- WILL OPEN **SATURDAY 24/02, morning**
- Open for 36 hours
- Make use of the non-assessed ones that I put up every week for practice
  - Also to know how to save, restart, submit, etc.
- **MATERIAL:** All we have covered till end of Week 6

- **LECTURE WEEK 7**

- **Exam preparation (well, sort of)**
- **Anatomy of a good answer...**
- Look at exam questions, look at what good and poor answers look like

# From Week 2



- Many processes share the memory
- But **where** in memory are they?
- Are they actually nicely arranged in contiguous blocks like in the picture?
- The **compiler** we use for our programs needs to know where in memory to put program data
  - But it does not know, the picture proves it – why?
  - Because it can not possibly know how many other 'things' are running and where in memory they are

# Memory Management Requirements

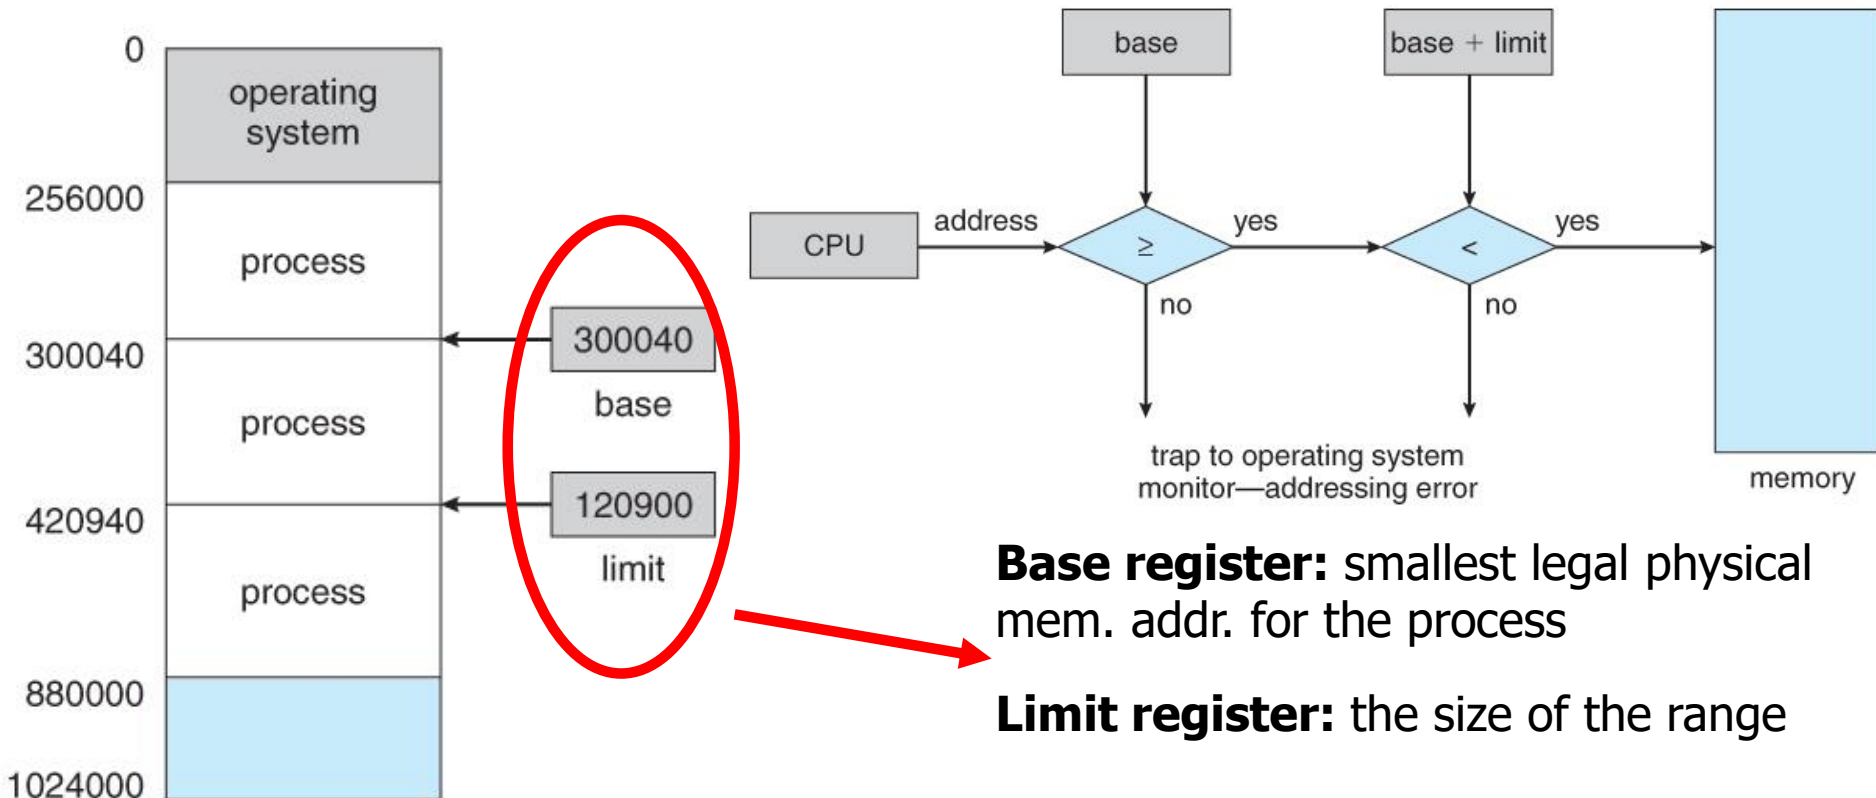
- **Why is memory management necessary?**
  - The computer memory is shared by (many) processes
  - But the compiler (and us) must produce code disregarding this fact – we can not know how many other processes are and where they are in memory
- Main **requirements** for memory management
  - **Relocation** (the ability to move the program from one place in memory to another)
  - **Protection & sharing**
  - Each process working in its own **address space** - (illusion)

# Req. 1: Relocation

- **Where** is the program in memory?
  - the programmer does not know, the compiler does not know
  - it may move around after it is loaded, will def. move around if you stop it and run it later
  - what if you run two instances of the same program? where will they be?
- We need to **be able to relocate a program in memory**
- **What** really are variables?
  - References to memory locations
- **How** is this simple instruction executed  $x = x + 1$ ?
  - In a series of LOAD and STORE assembly instructions

# Req. 2: Protection and Sharing

- No use of memory locations in another process
  - Accessing memory that the program is not allowed to causes a **trap** (check the term from week 2) – switch into kernel mode for OS to act
- But there are cases where we may want to allow several processes to access the same portion of memory
  - Inter-Process Communication (IPC), shared libraries, ...





# Addresses: 2 types (for now)

- **Logical**

- Reference to a memory location independent of the current use of memory
- Instructions issued by CPU reference logical addresses (pointers, arguments to LOAD/STORE, etc.)

- **Physical** or Absolute

- The actual location in memory
- Used by CPU
- Easy to remember: physical memory


the memory  
you paid  
for!!!

- Address translation

- **Change the address** : logical  physical


# Logical Address Space

- **Logical Address space** is the set of addresses available to a program
- Possible addresses
  - 16 bits = Not enough any more
  - 32-bits = 4GBytes
  - 64-bits = (A LOT!!!)
- Recall these different segments from Week 2?



# Physical (Real) Memory

- **The memory we paid for**
- Real memory is always less than logical address space
  - **Memory mapped I/O**
  - **OS data**
  - **Simply can not afford  $2^{64}$  memory**
- **Address space abstraction:**
  - Process has its own address space
  - Can you imagine how it would be without this abstraction?



**e.g.: 4GB of physical memory**

# An old Solution – Partition Memory


- **Fixed Partitioning**

- Divide memory into contiguous 'blocks' at boot time
- Partitions don't change – what happens if your program grows? (you run out of space)
- **Internal fragmentation** (we have reserved some memory for a process which we are not fully using)
- Hardware requirement: **base** register (loaded by OS at context switch)
  - **physical address = logical address + base register**

- **Dynamic Partitioning**

- Create partitions as programs are loaded
- **External fragmentation** (unused memory broken-up into small regions that are unusable), but no internal fragmentation
- Hardware requirement: **base and limit register**
  - **physical address = logical address + base register**
  - **protection:** if (physical address > (base+limit)) then **TRAP!!!!**

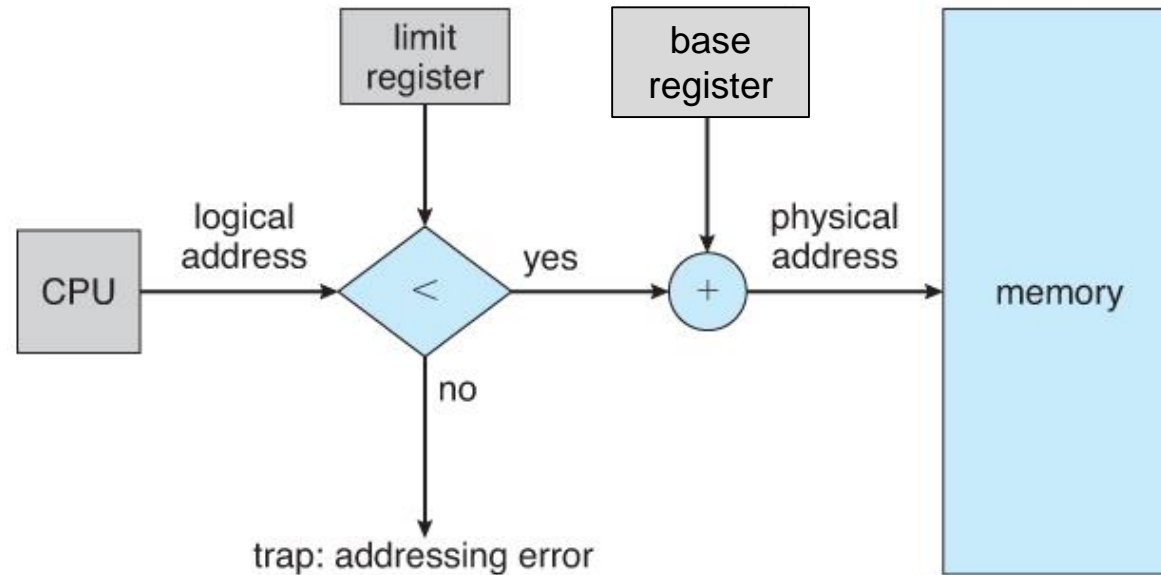
# Dynamic Partitioning Example



- **External Fragmentation**
- Memory external to all processes is fragmented
  - In this scenario, a new process requiring 7M of memory would not fit, even if the total free memory available is much over 7M
  - We would need to compact the memory to make the free space contiguous – this has a high overhead

Concept link: (fragmentation in disks)

# Protection and relocation in dynamic partitioning



- **protection** against user programs accessing areas that they should not
- programs can be **relocated** to different memory addresses as needed

Uses base and limit registers

- **base:** the 'starting point' of the process in physical memory
- **limit:** the 'size' of the process

# Key Concepts - Checkpoint

## 1. Relocation

- Code can be moved from one location in memory to another

## 2. Fragmentation

- A problem to avoid: unusable chunks of memory
- Internal and external

## 3. Logical versus physical addresses

- Logical: from the compiler, as seen by the CPU
- Physical: the memory we paid for

## 4. Address translation

- Change logical addresses to physical
- One method of relocation

# Better Solutions: Pages and Segments

- **Segmentation**

- Divide program into segments
  - Each segment is contiguous
  - Some external fragmentation

- **Paging**

- Divide memory into equal-size pages
- Load program into available pages
  - Pages do not need to be consecutive
- Some internal fragmentation

- **Virtual memory** (Week 8)

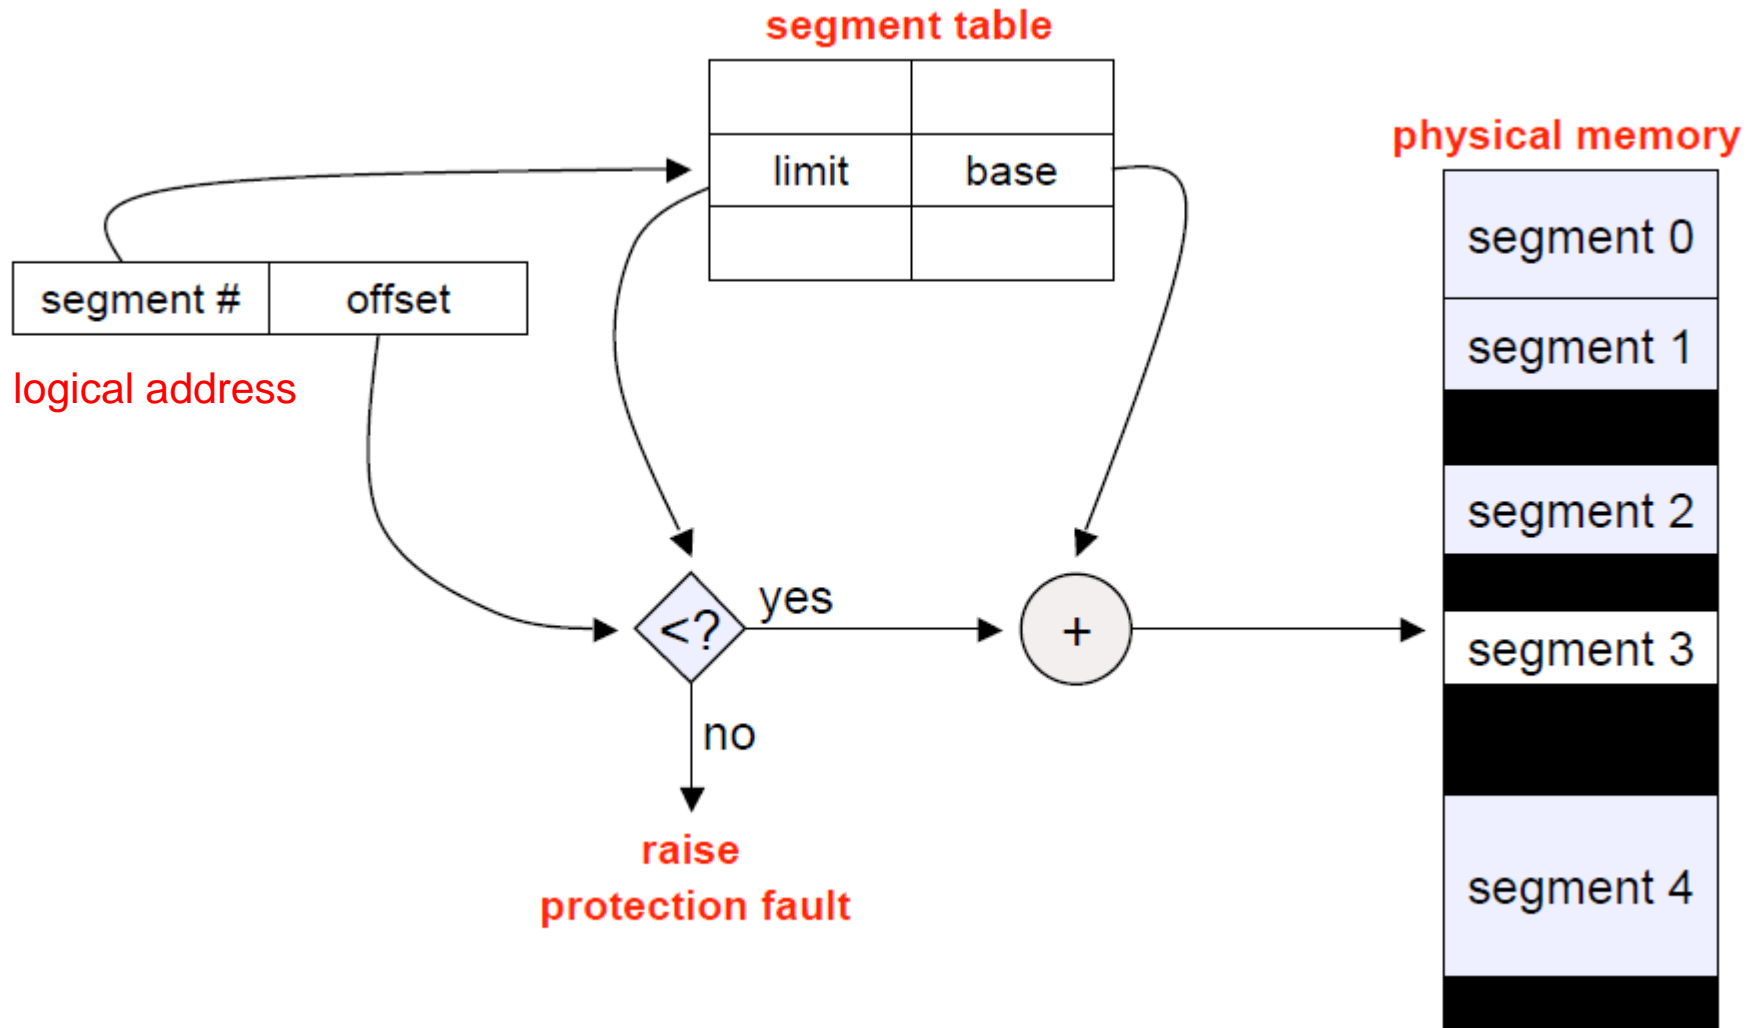
- “when our programs need more memory than we paid for...”



# Segmentation

- **A program can be subdivided into segments**
  - Segments may vary in length
  - There is a maximum segment length
- Why do we do that?
  - e.g. we want to separate the code from the data
  - **Code is read-only, data can be r+w**
  - Because code is read-only we can also safely share it if we want amongst many processes
- Addressing consists of two parts
  - a segment number and
  - an offset
- Segment table for a process

# Segment look-ups



# Segmentation Pros and Cons

- **Pros**

- Protection, as we can set segments for
  - Data: read and write
  - Code: read only
  - Shared
- No internal fragmentation

- **Cons**

- Segment contiguous in memory (not flexible)
- External fragmentation

# Looking ahead: Labs 6 & 7

- The next 2 labs use **SystemTap**
- **No more PHP ☺ ... but still vminstance ☹**
- Allows to look into the kernel, we need to be 'root' so we run them using vminstance in ITL
- **Lab 6 (week 8):** Not assessed, looking at CPU burst lengths (scheduling)
  - Helps you learn to use system tap and prepare for the assessed lab
- **Lab 7 (week 9):** Assessed, looking at memory management (paging, page faults, etc.)
- Good (?) news: not a single line of code to write
- It requires (allows) you to **link what you observe in the lab with what we say in the lectures**

# Paging: Frames and Pages

- Fixed size memory **frames** (refers to **physical address**)
- Process has memory **pages** (**same size as frames**) – refers to **logical address**
- We need to **allocate** a frame to each page that our process is using
  - any page (from any process) can be placed into any available frame
- Allows processes' physical memory to be **discontinuous**
- **Advantage** (in terms of efficient use of memory)
  - No external fragmentation
  - Internal fragmentation ok if page size (and so also frame size) is small

# Page size trade-off

Frame number	Main memory
0	A.0
1	A.1
2	A.2
3	A.3
4	D.0
5	D.1
6	D.2
7	C.0
8	C.1
9	C.2
10	C.3
11	D.3
12	D.4
13	
14	

- The example shows that we avoid external fragmentation and that we allow for processes to grow (in pages)
- **Page size is a trade-off**
  - **Too large:** we risk to have a fair amount of internal fragmentation
  - **Too small:** We end up with too many pages and we need to store data in memory for them (page table)
  - **Current typical size: 4KB** but trend is to increase it as memory becomes cheaper (we can afford to store more stuff)

# The Page Table

- Each process needs a map to show where its pages are (i.e. in which frame)
  - map **logical** → **physical address**
- We use a structure called a **Page Table for each process**
  - it is indexed by the page number

0	0
1	1
2	2
3	3

Process A  
page table

0	—
1	—
2	—

Process B  
page table

0	7
1	8
2	9
3	10

Process C  
page table

0	4
1	5
2	6
3	11
4	12

Process D  
page table

13
14

Free frame  
list

- **Protection:** each process only has access to its allocated frames
- **Sharing:** we can allow sharing on a page level

# Page Table Entries

- Each process has a page table:

Page Number	0	Page flags	Frame number
	1		
	2		
	3		
	4		
	5		
	6		

- Flags:**

- Is page in memory? (or it is on disk?)
- Has it changed?
- Protection: read only?
- Shared?



# Page Table Entries: Example (Intel x86)

Page Frame Number (Physical Page Number)	Free (OS)	0	L	D	A	PCD	PWT	U	W	P
31-12	11-9	8	7	6	5	4	3	2	1	0

P: Present (same as “valid” bit in other architectures)

W: Writeable

U: User accessible

PWT: Page write transparent: external cache write-through

PCD: Page cache disabled (page cannot be cached)


A: Accessed: page has been accessed recently

D: Dirty (PTE only): page has been modified recently

L: L=1⇒4MB page (directory only, for 2-level page tables).  
Bottom 22 bits of logical address serve as offset

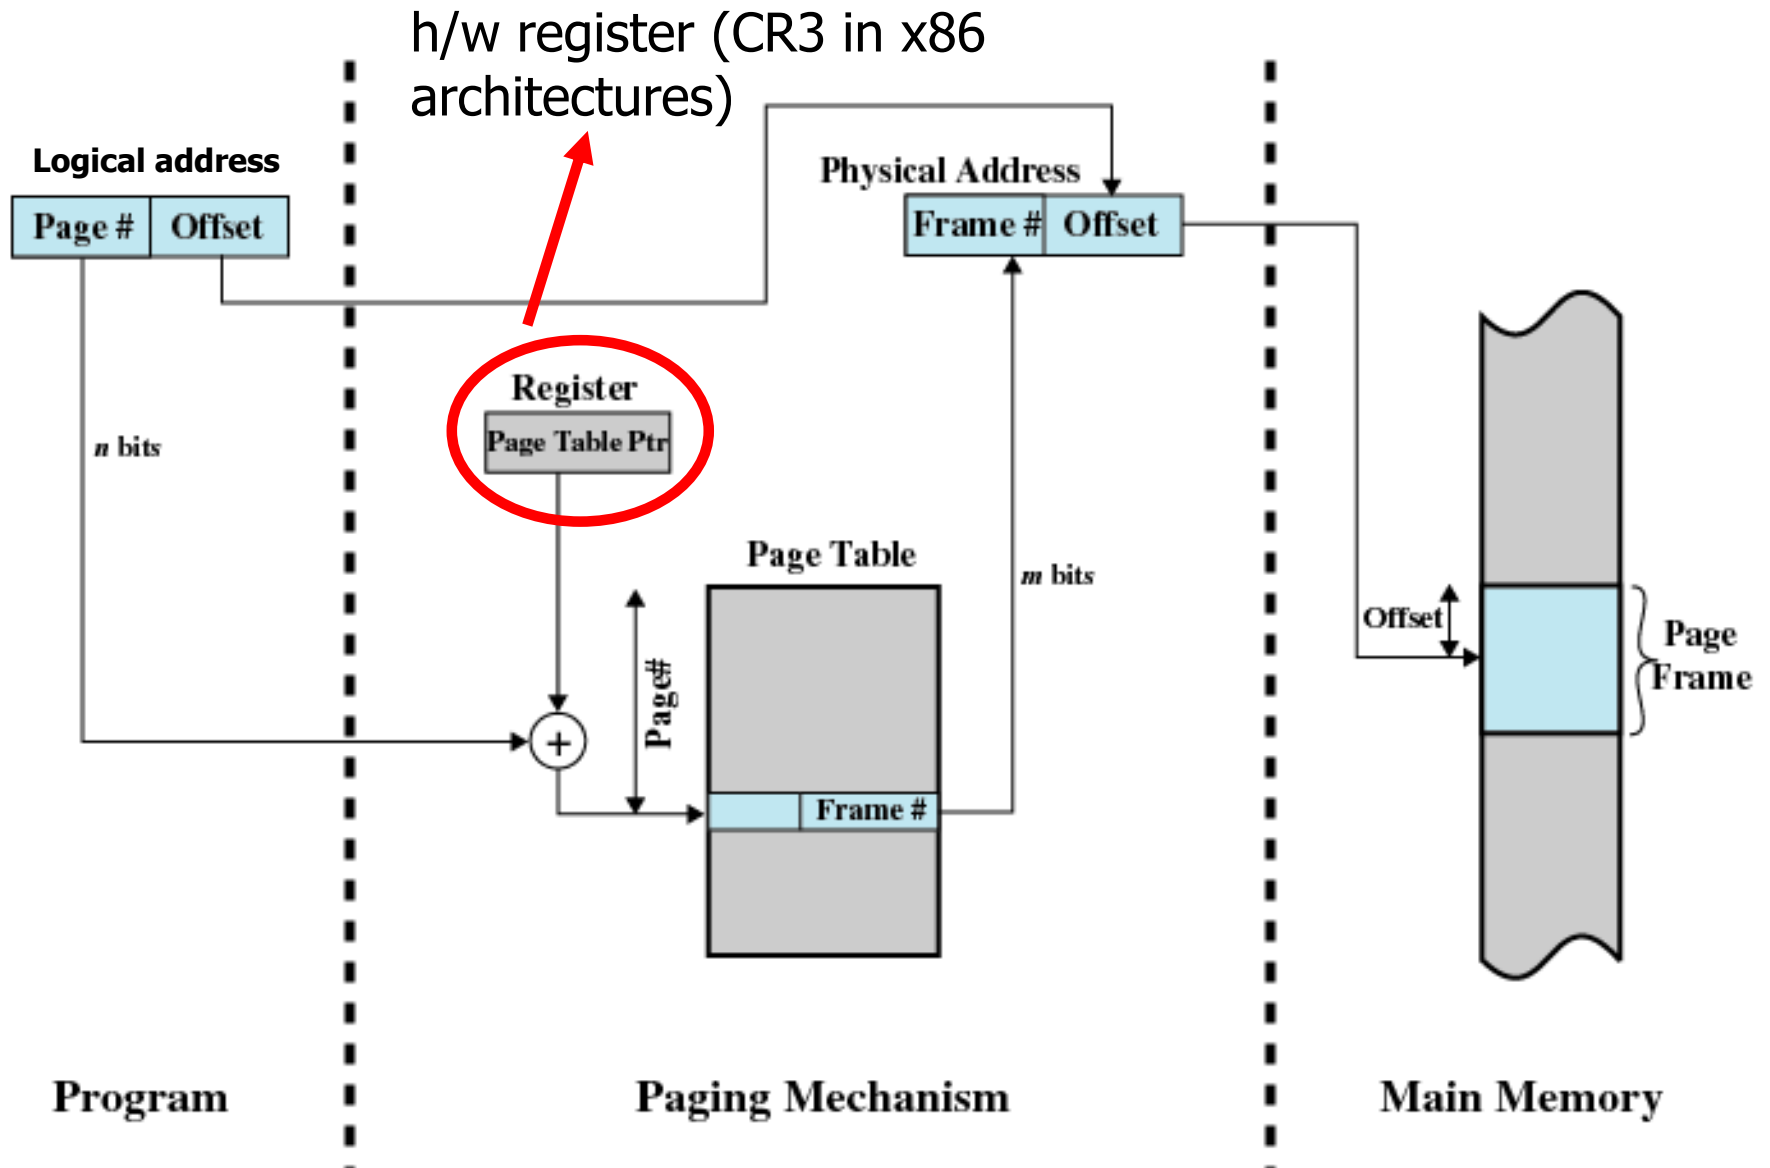
# Page Sizes and Logical Addresses

- A logical address for paging is split into:
  - Page number – **most significant bits**
  - Offset – **address within page**
- The actual split (e.g. 20/12) is architecture dependent



- From this we know:
  - **Page size:**  $2^{12} = 4\text{KB}$  (same as frame size)
  - **Number of pages:**  $2^{20} = 1 \text{ million pages}$
  - **Size of page table for each process, for 4bytes (=32 bits, as shown in previous slide) per Page Table Entry? 4MB**

# Address Translation I



# Examples

Can you figure out this example?



0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

logical memory

Page 0

Page 1

Page 2

Page 3

0	5
1	6
2	1
3	2

page table

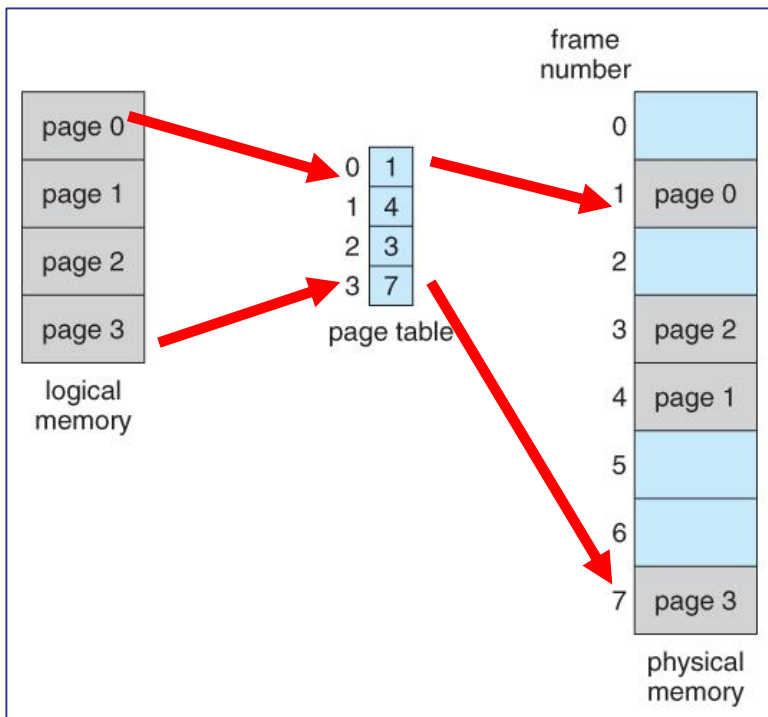
Frame 0

Frame 1

Frame 5

0	
4	i
8	j
	k
	l
	m
	n
	o
	p
12	
16	
20	a
	b
	c
	d
24	e
	f
	g
	h
28	


physical memory



# Example

What is great with  
16-bit processors?

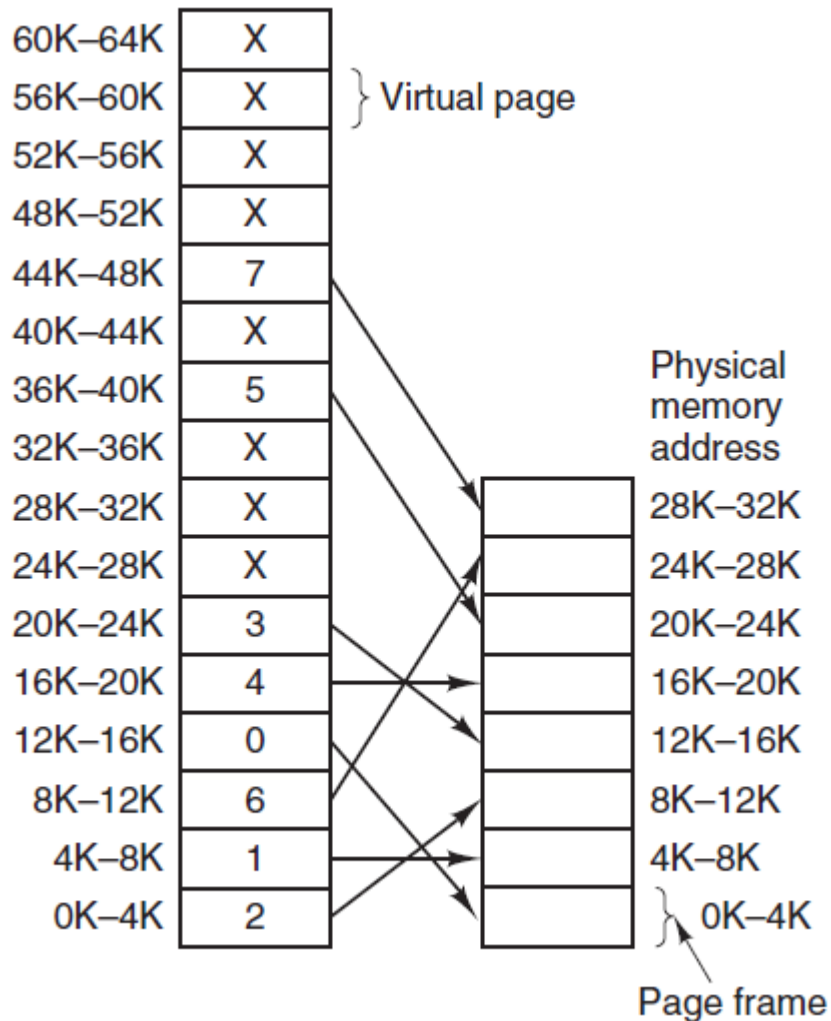
It's easy to follow  
translation  
examples in  
binary!!!



# Paging Advantages

- Eliminates **external fragmentation**
- Easy to **implement**
- Easy to model **protection**
- Easy to model **sharing**
- **Easy to allocate** physical memory
  - Allocate a frame from list of free frames
- Leads naturally to **Virtual Memory**
  - It is not necessary to have the whole program 'memory resident'
  - We can take pages that we don't need off main memory to a page file somewhere (on disk)

# Virtual Memory and Page Faults (link to next lecture)



- The **P** bit in the PTE of x86 keeps track of which pages are physically present in memory
- Pages marked with **X** are not memory resident, but kept on virtual memory (disk) – they have the **P** bit unset
- Requesting an unmapped page will cause the CPU to **trap** to the OS
- This trap is called a **Page Fault**
- **Week 8:** How we deal with Virtual Memory, page faults, strategies for caching, replacing pages, etc.

# Some problems with Paging

- **P1: Page table is large**
  - 32 bit addresses & 4 KByte pages → 4 Mbyte per page table (assuming 4 bytes / entry)
  - **For Each Process!**
  - **Much worse for 64bit addresses**
  - But... are all 1 million Page Table Entries (PTEs) needed?
    - **No, sparse tables**
- **P2: Memory access is slow** (in CPU terms...)
  - TWO memory look-ups for each memory access
  - One look-up into the page table, one more to access the data in memory
- **Solutions** to these & more, Week 8



# Summary

- CPU generates instructions in **logical space**
- Hardware sees **physical address space** (the memory we paid for)
- OS provides the illusion (abstraction) of an **address space owned by a process**
- **Memory Management** provides support for translation, relocation, protection, sharing, etc.
- Solutions need to consider issues such as efficient use of memory (**fragmentation**), **speed** (access to main memory is slow compared to CPU speed), ...